



CRIPAC

智能感知与计算研究中心
Center for Research on Intelligent Perception and Computing



中国科学院自动化研究所
Institute of Automation
Chinese Academy of Sciences

2019

“Deep Learning Lecture”

Lecture 1 : Introduction

Wang Liang

Center for Research on Intelligent Perception and Computing (CRIPAC)
National Laboratory of Pattern Recognition (NLPR)
Institute of Automation, Chinese Academy of Science (CASIA)

Outline

1/ Course Information

2/ What is Deep Learning

3/ Applications of Deep Learning

4/ Future Directions

5/ Project and homework guidance

Course Goals

- Give you an opportunity for you to explore an interesting multivariate analysis problem of your choice in the context of a real-world data set
- How deep learning works
- How to frame tasks into learning problems
- How to use toolkits to implement designed models
- When and why specific deep learning techniques work for specific problems

Course Prerequisites

- Calculus and Linear Algebra
- Probability and Statistics
- Machine Learning
- Advanced programming language (like Python and C)
- Time and patience

Course Overview

1. Introduction
2. Mathematical Basics
3. Feedforward Network
4. Convolutional Neural Network
5. Recurrent Neural Network
6. Regularization and Optimization
7. Deep Generative Models
8. Reinforcement Learning
9. Attention and Memory
10. RBM, DBN and DBM
11. Graphical Model
12. DL Application
13. Project Presentation

Course Logistics

- Course: Theory (3 hours per week) + course project + homework
- The location of course room: 教1-101
- Final grade = 50% final exam + 25% course project + 25% homework (paper reading)
- The course project contains a technical report and each team is composed of 5-10 members.
- The course project asks for volunteer presentation, otherwise it will be decided by random sampling

Course Logistics

- We strongly encourage students to form study groups. Students may discuss and work on project problems in groups.
- However, each student must write down the solutions independently, and without referring to written notes from the joint session. In other words, each student must understand the solution well enough in order to reconstruct it by him/herself.
- In addition, each student should write on the problem set the bunch of people with whom s/he collaborated.

Course Logistics

Projects will be evaluated based on:

1. The technical quality of the work. (I.e., Does the technical material make sense? Are the things tried reasonable? Are the proposed algorithms or applications clever and interesting? Do the authors convey novel insight about the problem and/or algorithms?)
2. The completeness and novelty of the work, and the clarity of the write-up. (Spare enough time for the write-up since it may be harder than you imagine)

Course Logistics

Your final report is expected to be a 4~8 page report. You should submit both an electronic and a hardcopy version for your final report. It should roughly have the following format:

- **Introduction - Motivation**
- **Problem definition**
- **Proposed method**
 - Intuition - why should it be better than other methods?
 - Description of its algorithms
- **Experiments**
 - Description of your testbed; list of questions your experiments are designed to answer
 - Details of the experiments; observations
- **Conclusions**

Course Logistics

- Write final reports in the form of CVPR papers:
http://cvpr2019.thecvf.com/submission/main_conference/author_guidelines (using Latex, which is very useful for your writing papers in later research life)
- Provide corresponding codes, and a “readme.txt” file to illustrate how to run the codes

People



Professor: Liang Wang
Email: wangliang@nlpr.ia.ac.cn



Associate Professor: Wei Wang
Email: wangwei@nlpr.ia.ac.cn



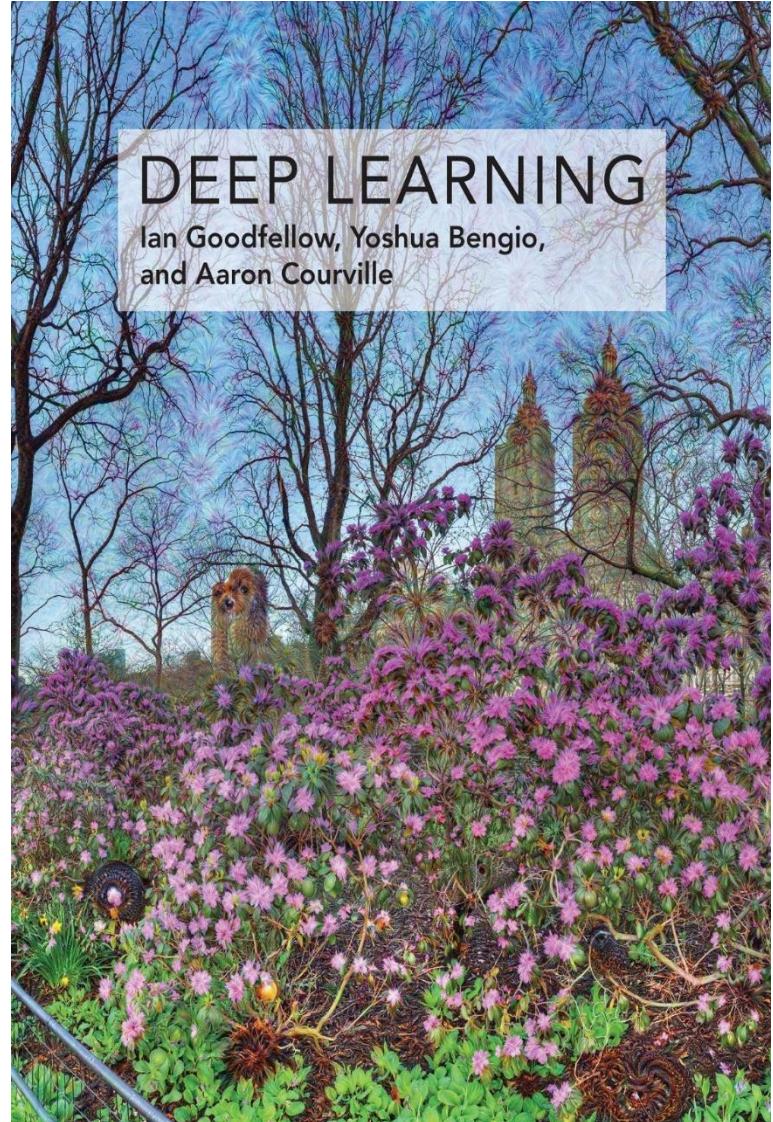
Teaching Assistant: Kai Niu
Email : kai.niu@cripac.ia.ac.cn



Teaching Assistant: Chenyang Si
Email : chenyang.si@cripac.ia.ac.cn

Reference book

- http://www.deeplearningbook.org/lecture_slides.html
- https://github.com/goodfeli/dlbook_exercises



Some Resources: Programming

- <http://scikit-learn.org/stable/> (Scikit Learn: a Python library for machine learning)
- <http://ubcmatlabguide.github.io/> (YAGTOM: Yet Another Guide TO Matlab)
- <https://github.com/saitjr/C-Toturials> (Learning C)

Some Resources: Deep Learning Tools

| Name | Language | Link | Note |
|----------------|----------|---|---|
| Pylearn2 | Python | http://deeplearning.net/software/pylearn2/ | A machine learning library built on Theano |
| Theano | Python | http://deeplearning.net/software/theano/ | A python deep learning library |
| Caffe | C++ | http://caffe.berkeleyvision.org/ | A deep learning framework by Berkeley |
| Torch | Lua | http://torch.ch/ | An open source machine learning framework |
| Overfeat | Lua | http://cilvr.nyu.edu/doku.php?id=code:start | A convolutional network image processor |
| Deeplearning4j | Java | http://deeplearning4j.org/ | A commercial grade deep learning library |
| Word2vec | C | https://code.google.com/p/word2vec/ | Word embedding framework |
| GloVe | C | http://nlp.stanford.edu/projects/glove/ | Word embedding framework |
| Doc2vec | C | https://radimrehurek.com/gensim/models/doc2vec.html | Language model for paragraphs and documents |
| StanfordNLP | Java | http://nlp.stanford.edu/ | A deep learning-based NLP package |
| TensorFlow | Python | http://www.tensorflow.org | A deep learning based python library |
| PyTorch | Python | https://pytorch.org | A deep learning based python library |

Some Resources: Conferences

A very good project will comprise a **publishable or nearly-publishable** piece of work. So, for inspiration, you might also look at some recent machine learning and artificial intelligence research papers.

Main machine learning and artificial intelligence conferences are ICML, NeurIPS, CVPR, AAAI and IJCAI. You can find papers from the following websites:

<https://icml.cc/Conferences/2018/Schedule> (ICML-2018)

<https://nips.cc/Conferences/2018/Schedule> (NeurIPS-2018)

<http://openaccess.thecvf.com/CVPR2018.py> (CVPR-2018)

<https://www.aaai.org/Library/AAAI/aaai18contents.php> (AAAI-2018)

<https://ijcai-17.org/accepted-papers.html> (IJCAI-2017)

Outline

1 Course Information

2 What is Deep Learning

3 Applications of Deep Learning

4 Future Directions

5 Project and homework guidance

Deep Learning

ARTIFICIAL INTELLIGENCE

Any technique that enables computers to mimic human behavior



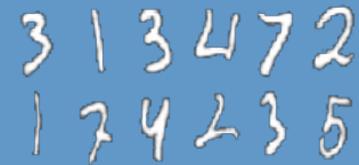
MACHINE LEARNING

Ability to learn without explicitly being programmed



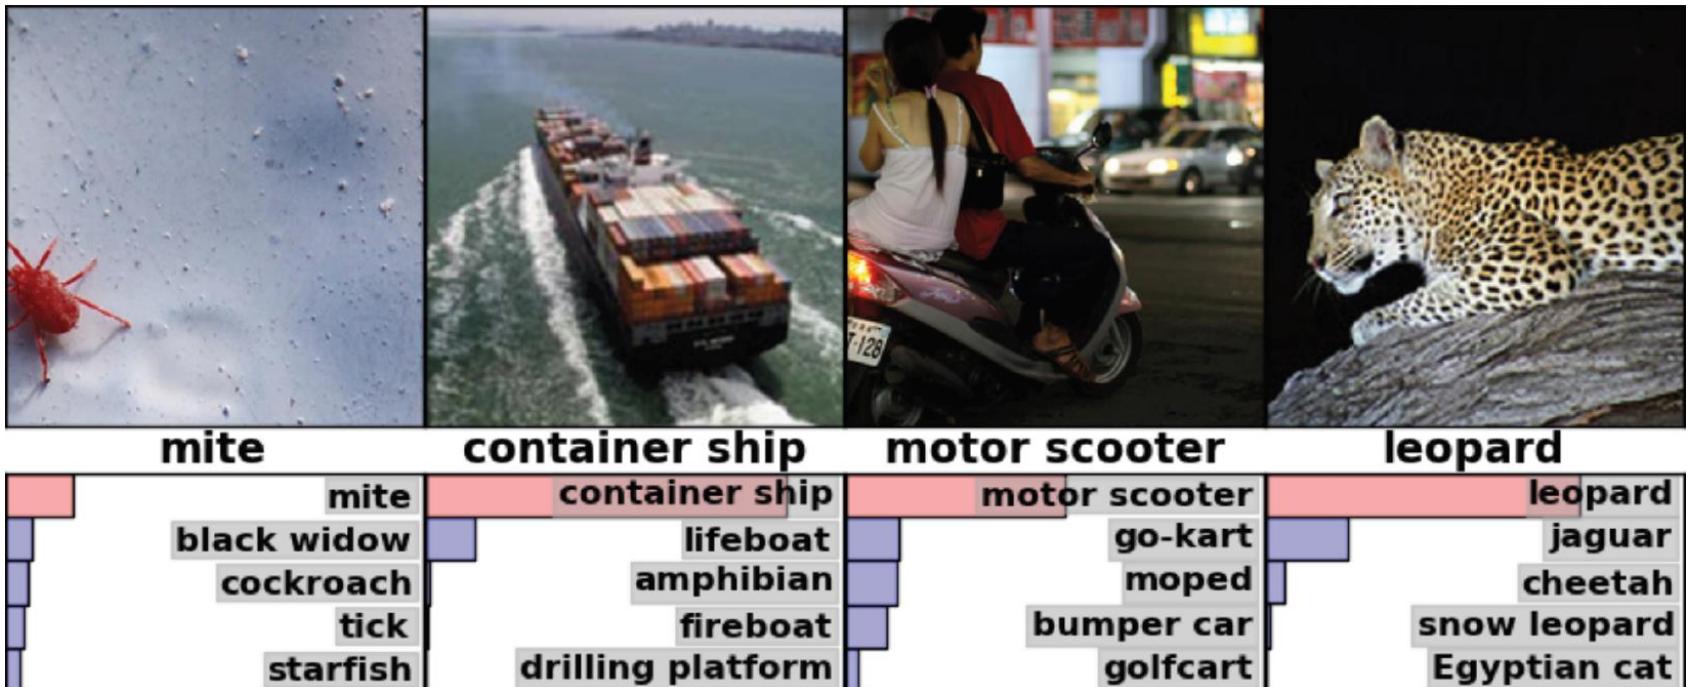
DEEP LEARNING

Learn underlying features in data using neural networks



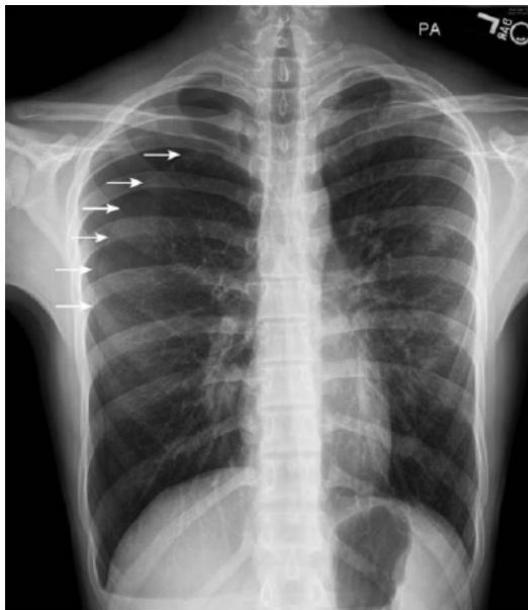
Deep Learning Success for Vision

Image Recognition



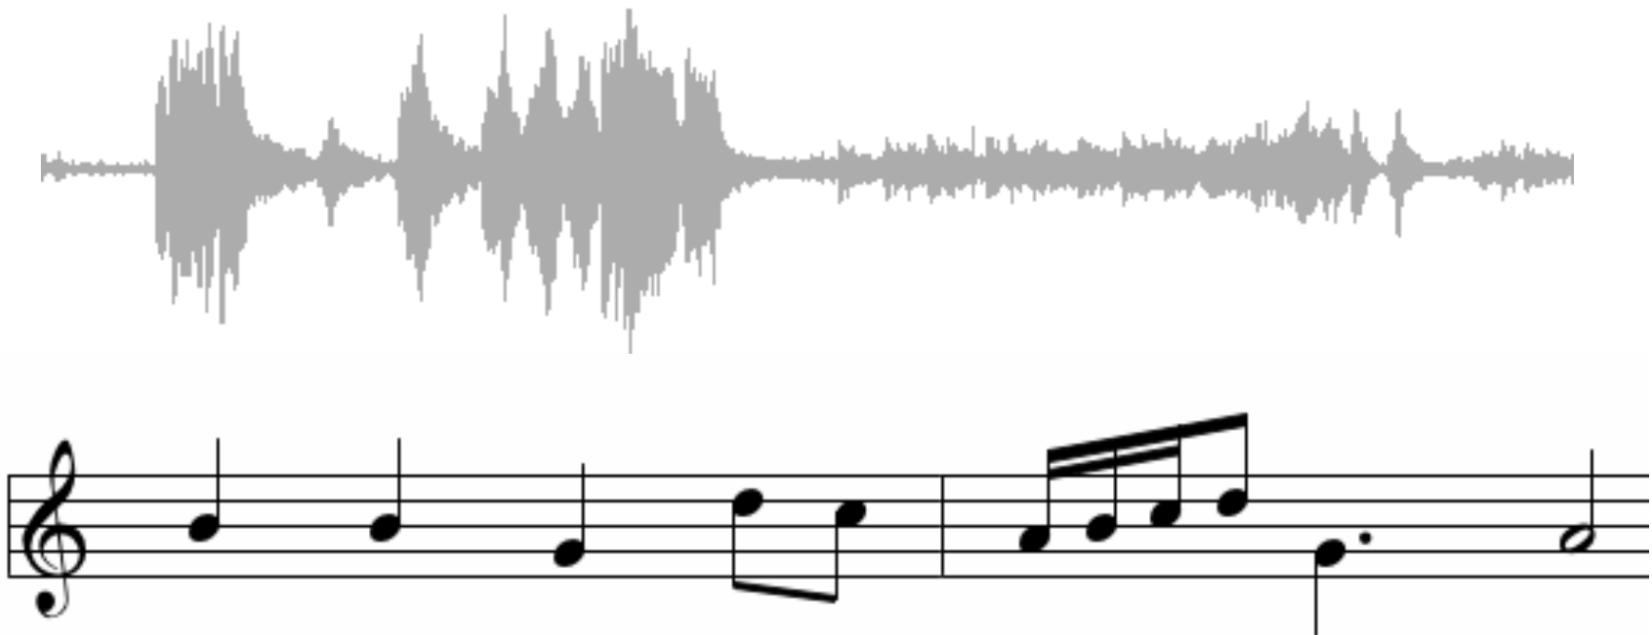
Deep Learning Success for Vision

Detect pneumothorax in real X-Ray scans



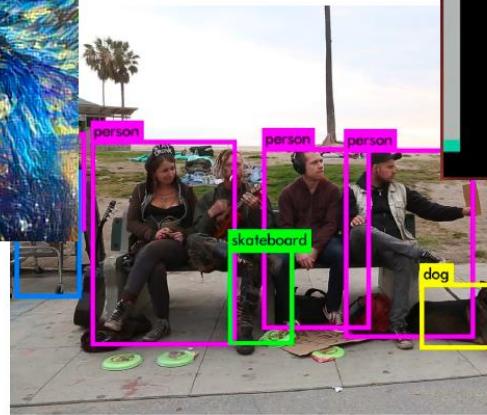
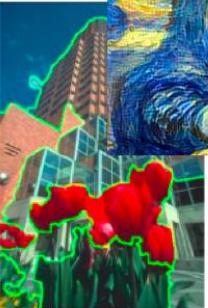
Deep Learning Success for Audio

Music Generation

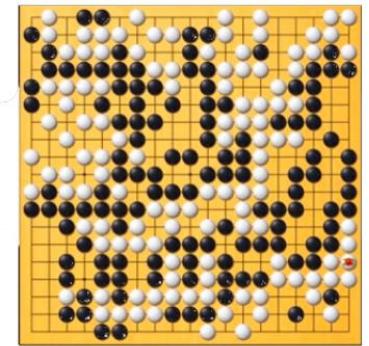


Deep Learning Success

So many more ...

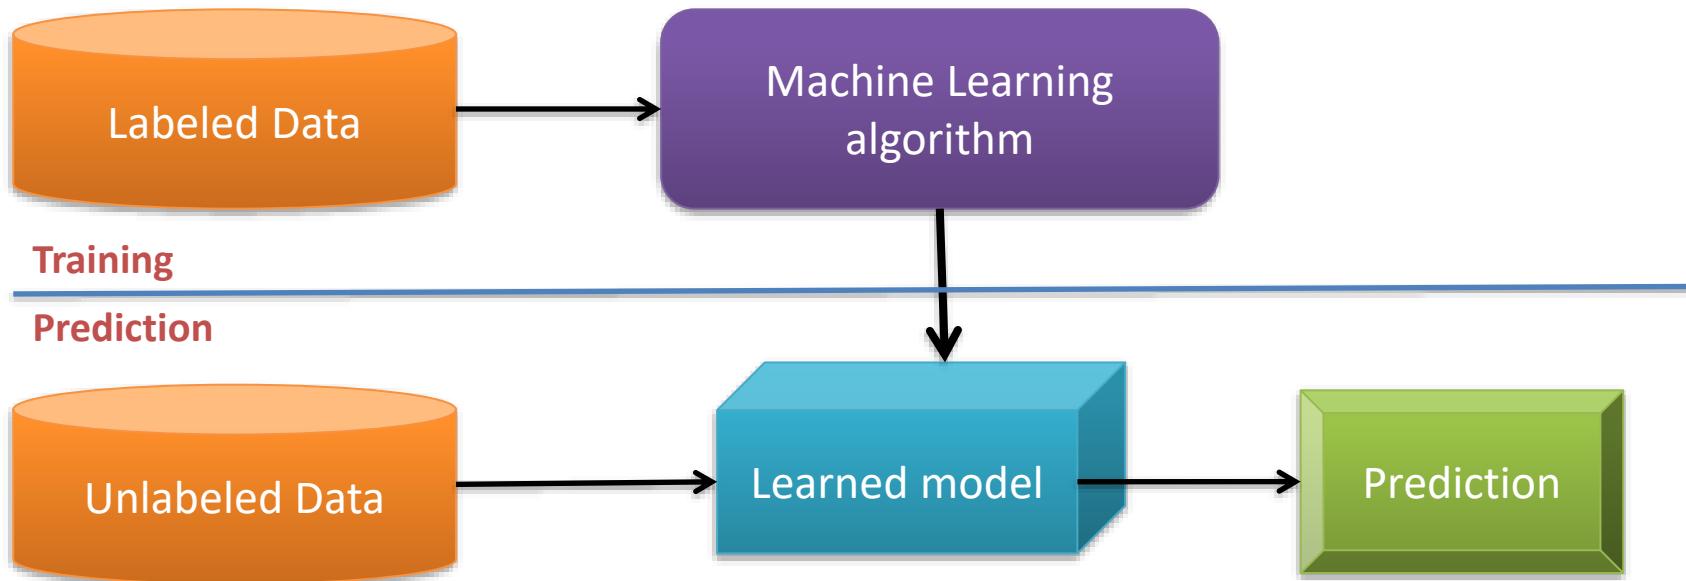


AlphaGo
ZERO



Machine Learning Basics

Machine learning is a field of computer science that gives computers the ability to **learn without being explicitly programmed**



Methods that can learn from and make predictions on data

Types of Learning

Supervised: Learning with a **labeled training** set

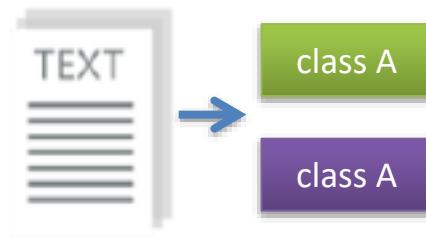
Example: email *classification* with already labeled emails

Unsupervised: Discover **patterns** in **unlabeled** data

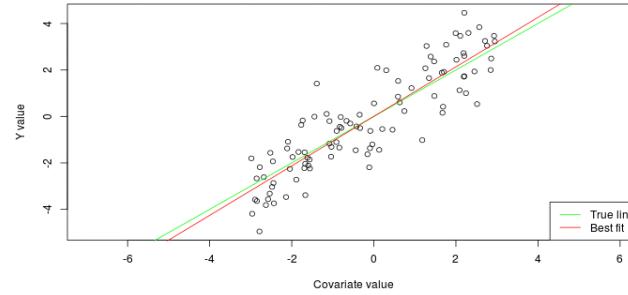
Example: *cluster* similar documents based on text

Reinforcement learning: learn to **act** based on **feedback/reward**

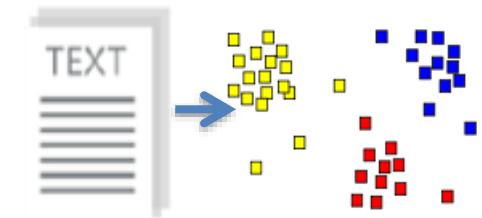
Example: learn to play Go, reward: *win or lose*



Classification



Regression



Clustering

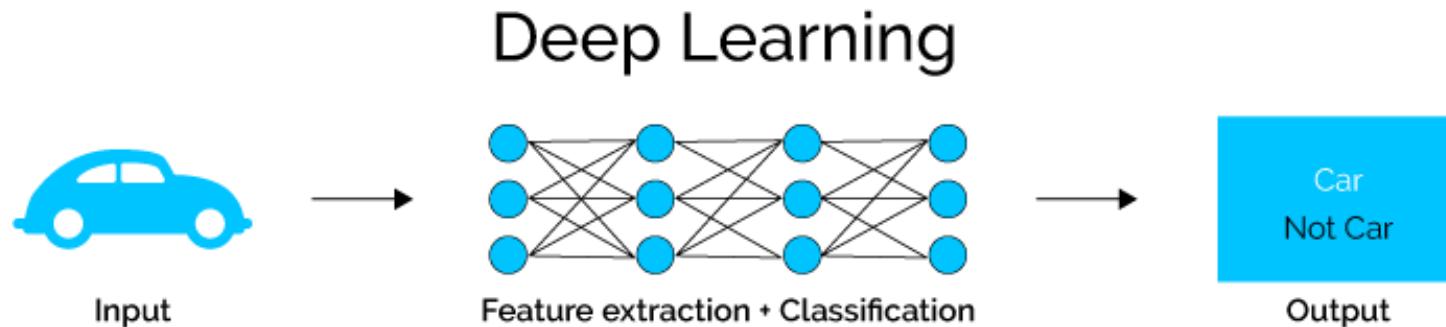
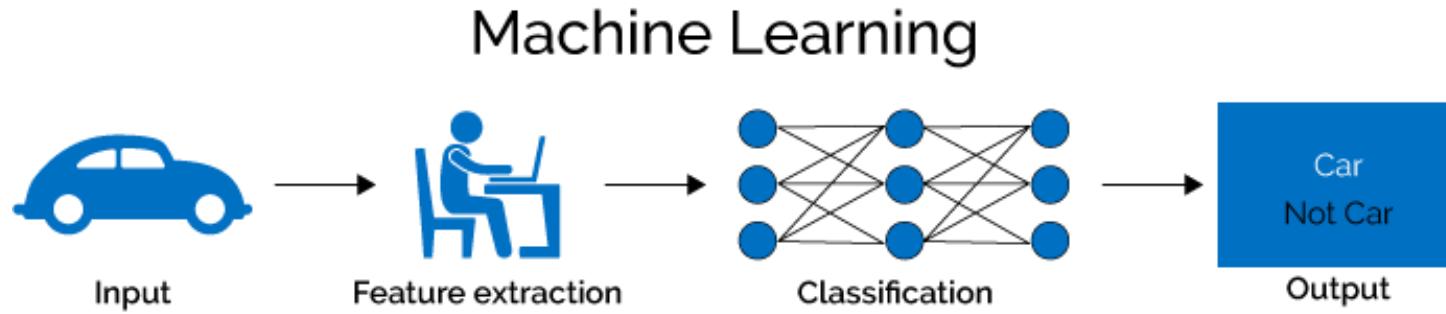
Anomaly Detection
Sequence labeling

...

ML vs. Deep Learning

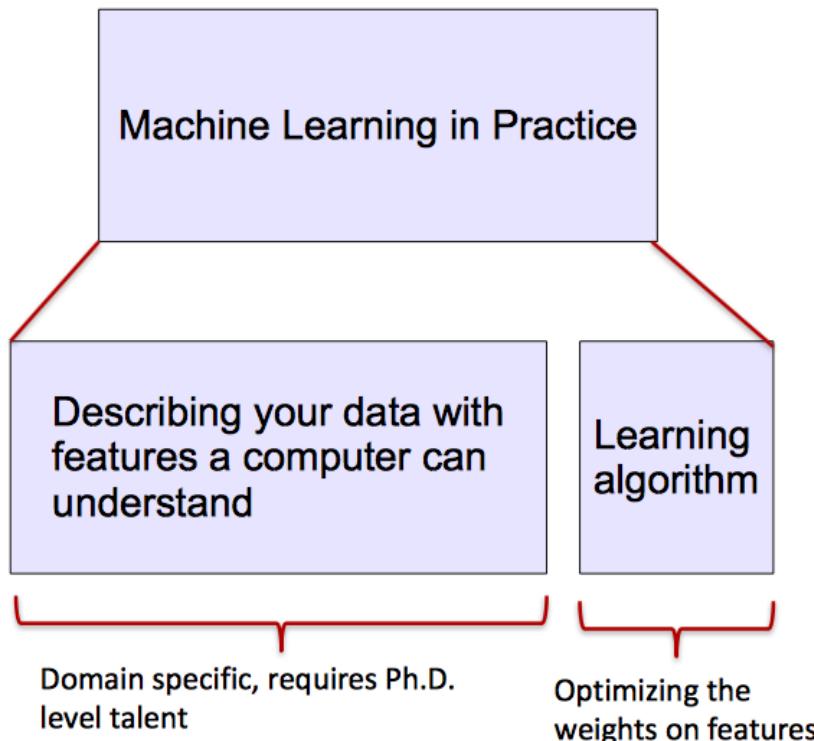
A machine learning subfield of learning **representations** of data.
Exceptional effective at **learning patterns**.

Deep learning algorithms attempt to learn (multiple levels of) representation by using a **hierarchy of multiple layers**
If you provide the system **tons of information**, it begins to understand it and respond in useful ways.



ML vs. Deep Learning

Most machine learning methods work well because of **human-designed representations** and **input features**
ML becomes just **optimizing weights** to best make a final prediction



| Feature | NER |
|----------------------------------|--------|
| Current Word | ✓ |
| Previous Word | ✓ |
| Next Word | ✓ |
| Current Word Character n-gram | all |
| Current POS Tag | ✓ |
| Surrounding POS Tag Sequence | ✓ |
| Current Word Shape | ✓ |
| Surrounding Word Shape Sequence | ✓ |
| Presence of Word in Left Window | size 4 |
| Presence of Word in Right Window | size 4 |

Deep Learning Today

- Advancement in **Speech Recognition** in the last 2 years
 - A few long-standing performance records were broken with deep learning methods
 - Microsoft and Google have both deployed DL-based speech recognition systems in their products
- Advancement in **Computer Vision**
 - Feature engineering is the bread-and-butter of a large portion of the CV community, which creates some resistance to feature learning
 - But the record holders on ImageNet and Semantic Segmentation are convolutional nets
- Advancement in **Natural Language Processing**
 - Fine-grained sentiment analysis, syntactic parsing
 - Language model, machine translation, question answering

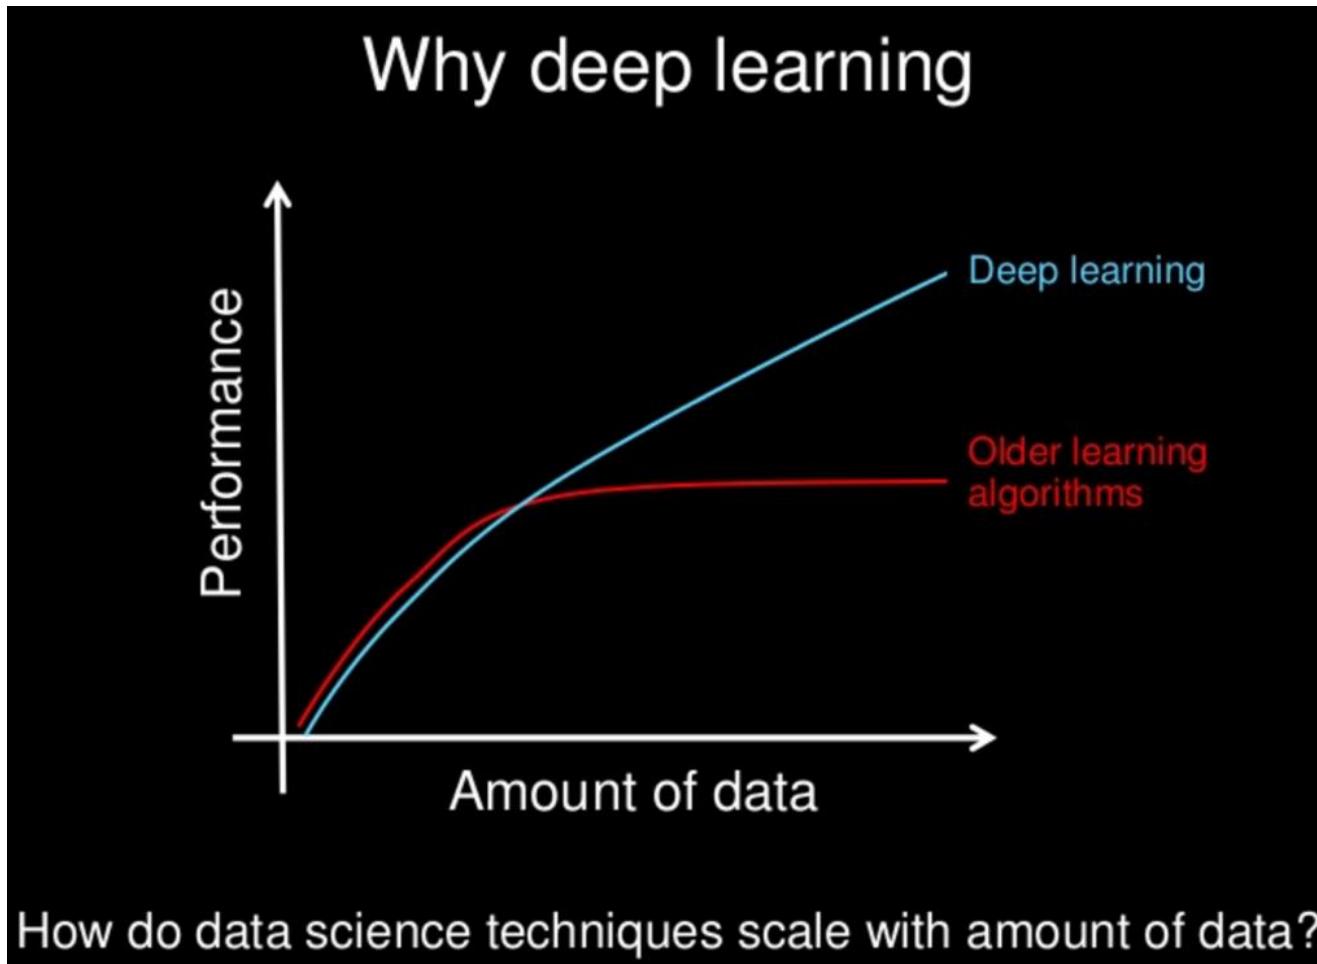
Deep Learning Today



In "Nature" 27 January 2016:

- "DeepMind's program AlphaGo beat Fan Hui, the European Go champion, five times out of five in tournament conditions..."
- "AlphaGo was not preprogrammed to play Go: rather, it learned using a general-purpose algorithm that allowed it to interpret the game's patterns."
- "...AlphaGo program applied **deep learning** in neural networks (convolutional NN) — brain-inspired programs in which connections between layers of simulated neurons are strengthened through examples and experience."

Why is Deep Learning Important



Deep Learning is finally enabling us to cross that line in places we weren't able to before.

Big Year for AI



1936: Turing Machine from Alan Turing



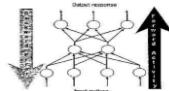
1946: Birth of ENIAC



1956: Birth of AI



1966: Turing Award



1986: The BP Algorithm



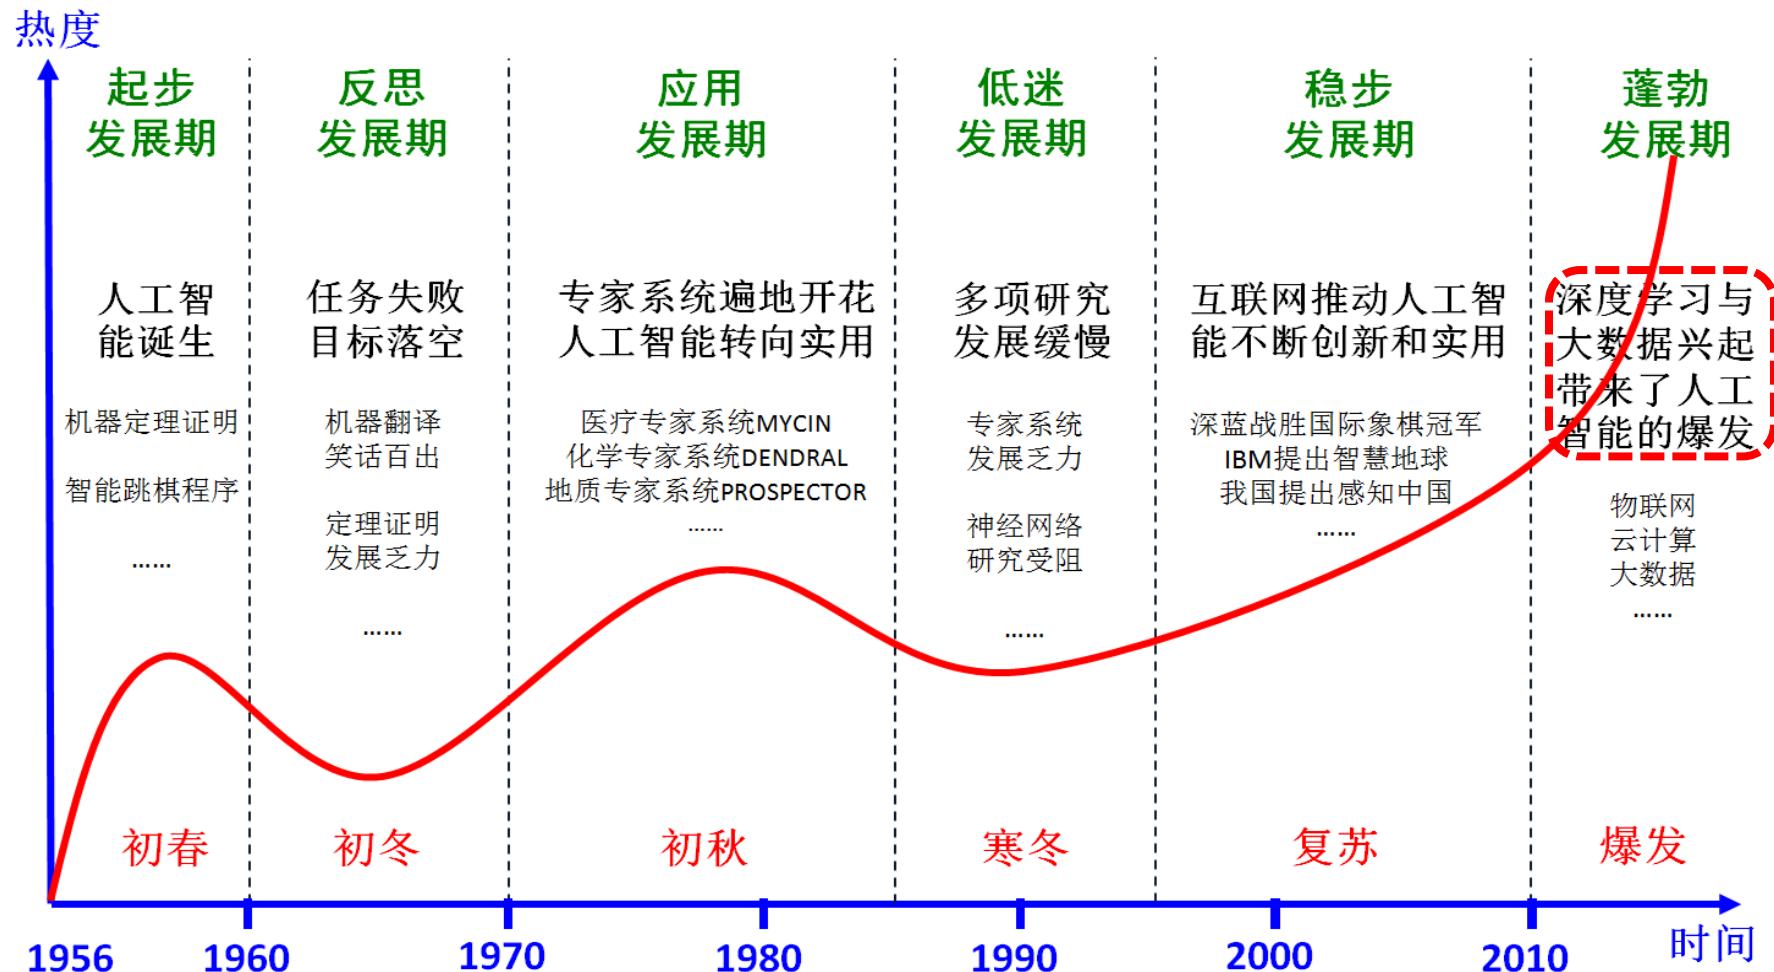
2006: DNN, Deep Learning



AlphaGo

2016: AlphaGo

The History of AI



Deep Neural Networks (DNN)

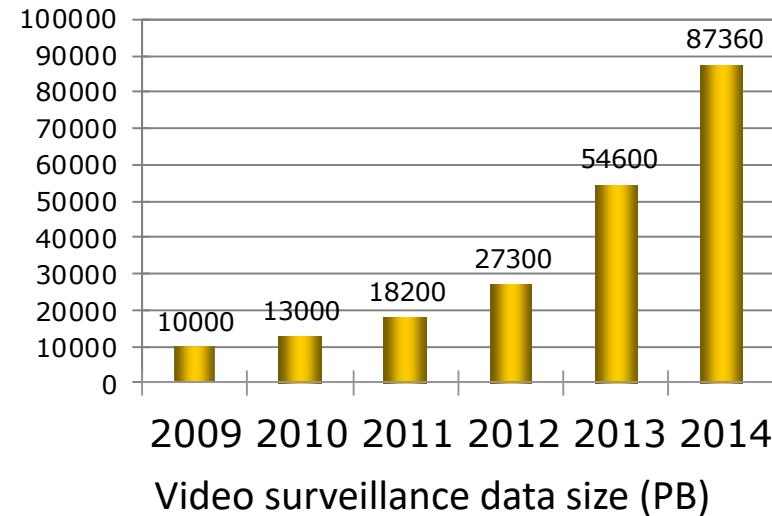
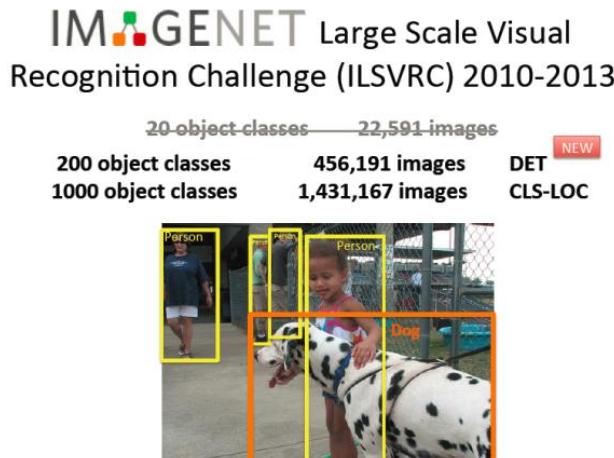
- **Originate from:**
 - 1962 – simple/complex cell, *Hubel and Wiesel*
 - 1970 – efficient error back propagation, *Linnainmaa*
 - 1979 – deep neocognitron, convolution, *Fukushima*
 - 1987 – autoencoder, *Ballard*
 - 1989 – convolutional neural networks (CNN), *Lecun*
 - 1991 – deep recurrent neural networks (RNN), *Schmidhuber*
 - 1997 – long short-term memory (LSTM), *Schmidhuber*
- **Two drawbacks:**

Large numbers of parameters → High computational cost

Small training set → Over-fitting problem

Two Recent Developments

Big Data



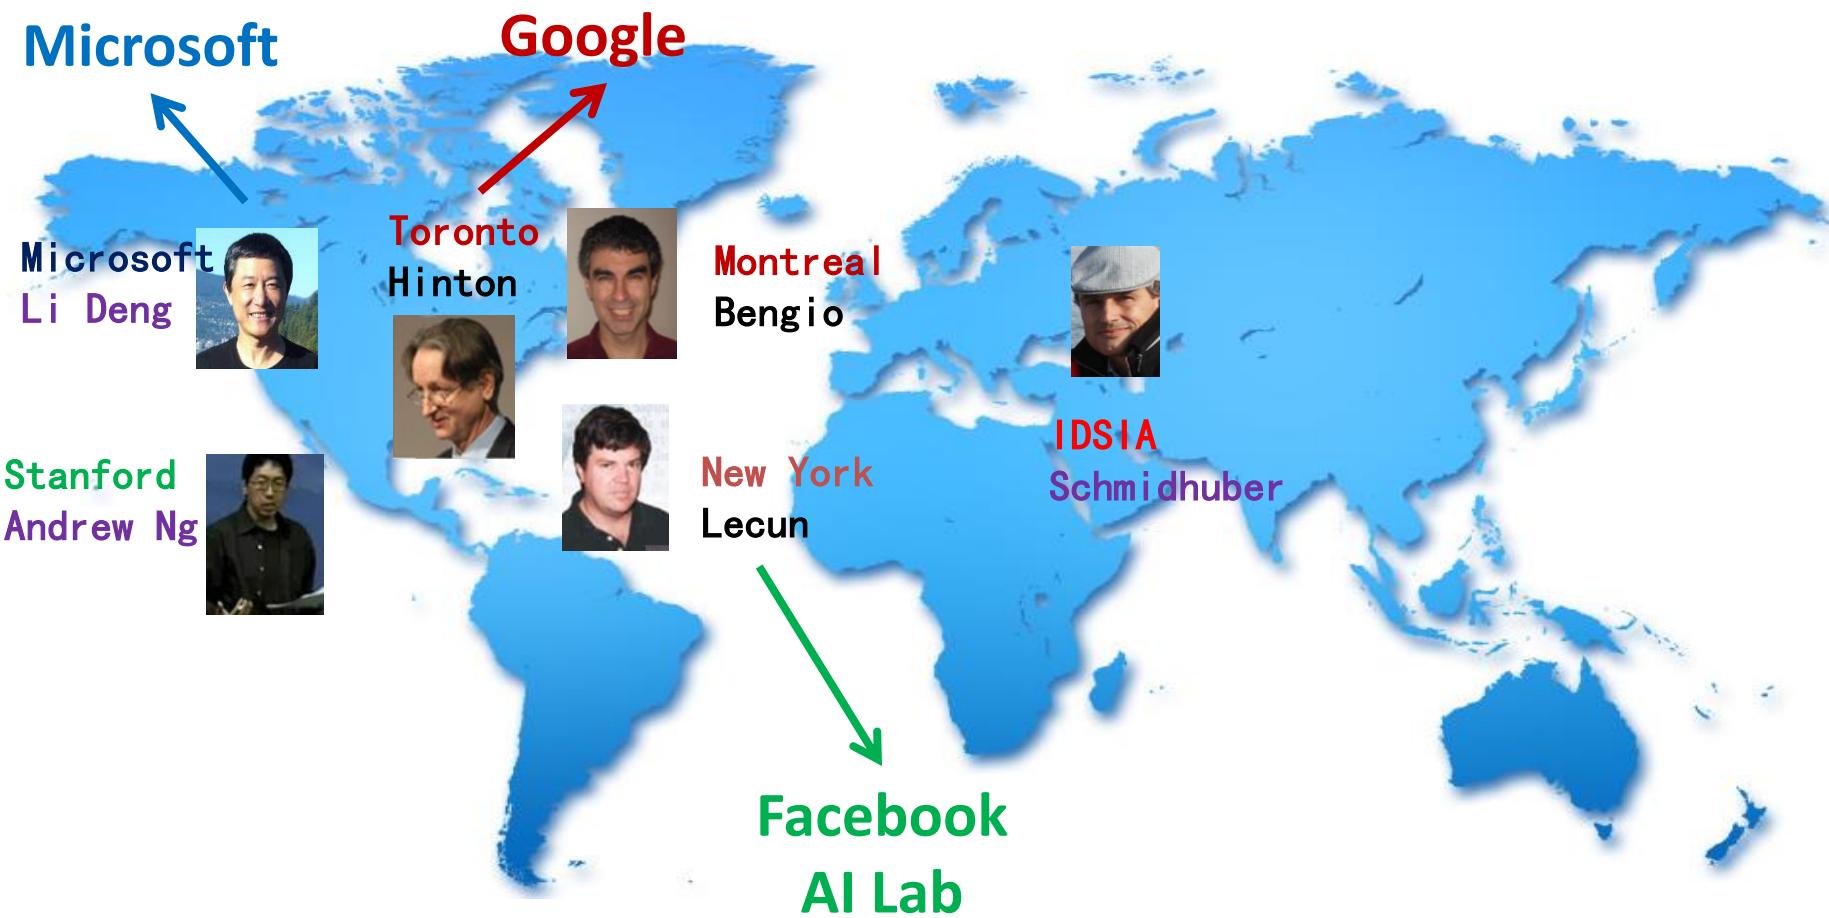
Cheap Computation



| Features | Tesla K40 | Tesla K20X | Tesla K20 |
|--|-----------------|-------------|----------------|
| Number and Type of GPU | 1 Kepler GK110B | | 1 Kepler GK110 |
| Peak double precision floating point performance | 1.43 Tflops | 1.31 Tflops | 1.17 Tflops |
| Peak single precision floating point performance | 4.29 Tflops | 3.95 Tflops | 3.52 Tflops |
| Memory bandwidth (ECC off) | 288 GB/sec | 250 GB/sec | 208 GB/sec |
| Memory size (GDDR5) | 12 GB | 6 GB | 5 GB |
| CUDA cores | 2880 | 2688 | 2496 |

DNN can thus be fitted efficiently

Deep Learning Pioneers (2006-)



They lead deep learning through three major stages!

Deep Learning: The Resurgence of DNN

Breakthrough in 2006

Reducing the Dimensionality of Data with Neural Networks

G. E. Hinton* and R. R. Salakhutdinov

High-dimensional data can be converted to low-dimensional codes by training a multilayer neural network with a small central layer to reconstruct high-dimensional input vectors. Gradient descent can be used for fine-tuning the weights in such "autoencoder" networks, but this works well only if the initial weights are close to a good solution. We describe an effective way of initializing the weights that allows deep autoencoder networks to learn low-dimensional codes that work much better than principal components analysis as a tool to reduce the dimensionality of data.

Dimensionality reduction facilitates the classification, visualization, communication, and storage of high-dimensional data. A simple and widely used method is principal components analysis (PCA), which finds the directions of greatest variance in the data set and represents each data point by its coordinates along each of these directions. We describe a nonlinear generalization of PCA that uses an adaptive, multilayer "encoder" network

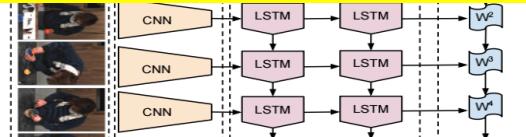
2006 VOL 313 SCIENCE www.sciencemag.org

2006

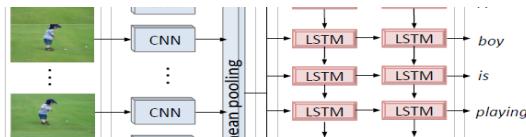
ImageNet: 74% vs. 85%



RNN for sequence analysis



Activity recognition, CVPR2015



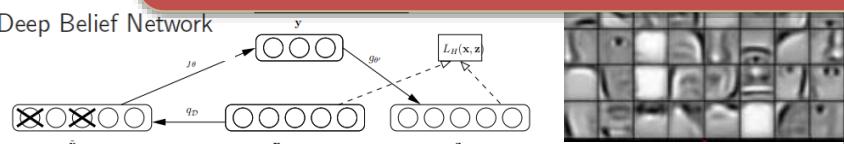
Video caption, CVPR2015

Representation learning



Deep Learning promotes the fast development of various pattern recognition areas

Deep Belief Network

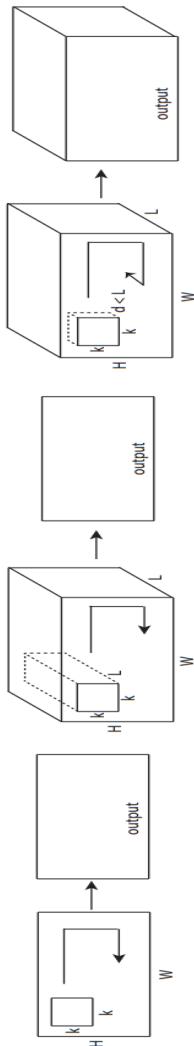


CNN for visual tasks

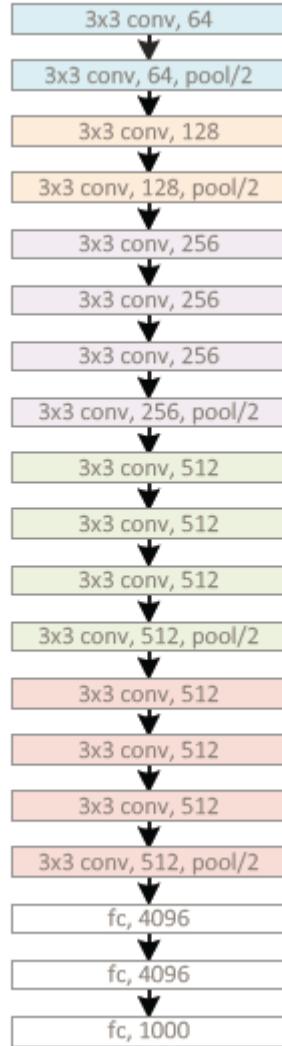


RCNN for detection, CVPR2014

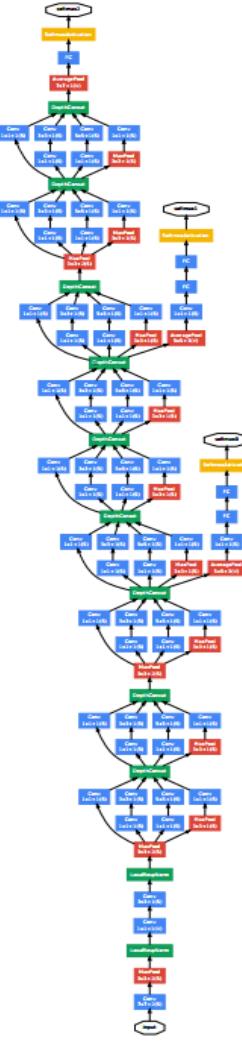
Deep Neural Networks



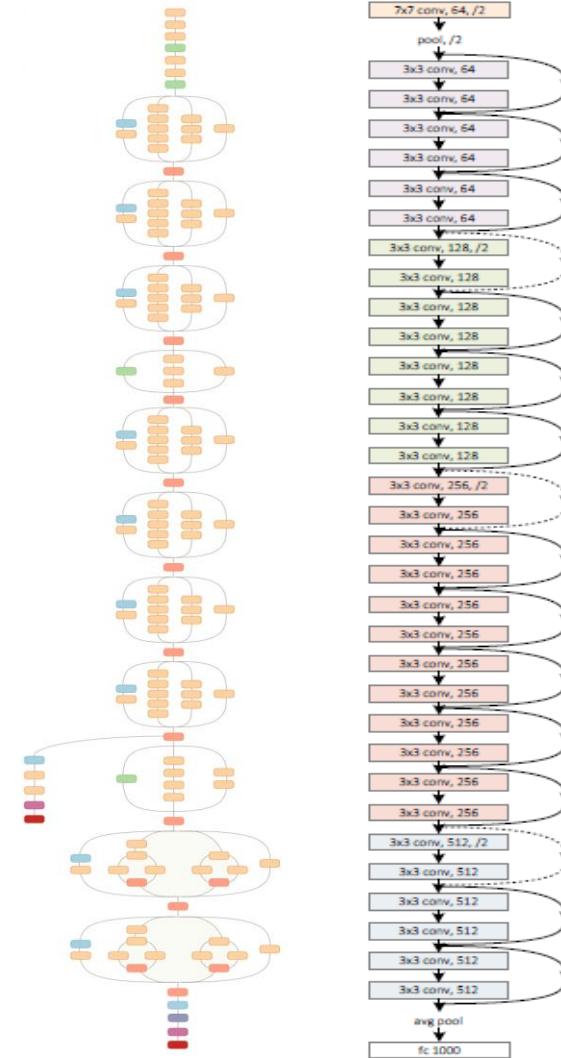
C3D



VGG-19



GoogleNet

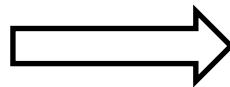
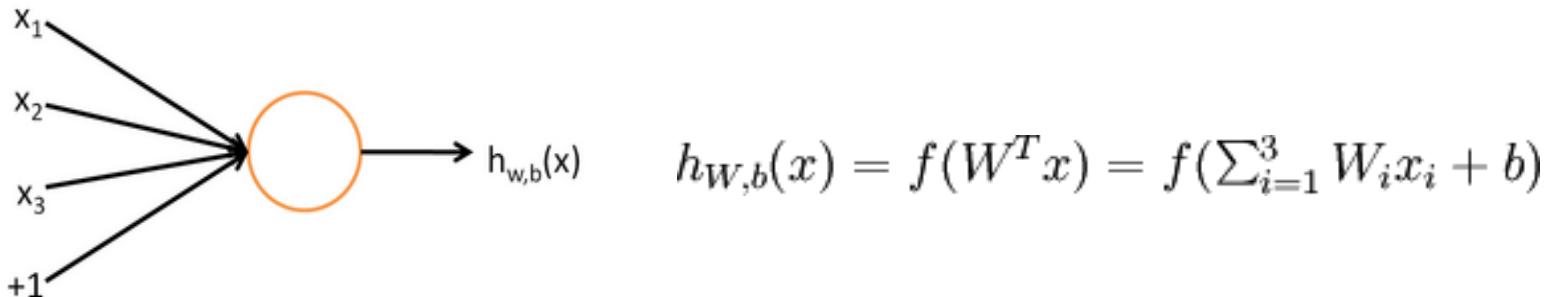


Inception-3

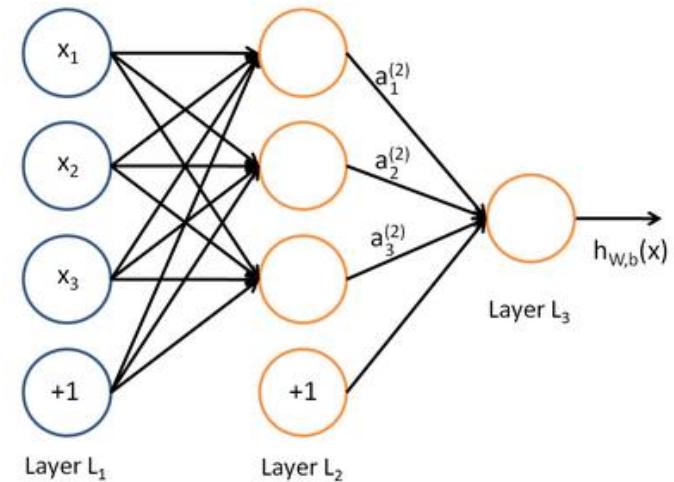
Residual

DNN: Full Connection

- **Each unit fully connects to all previous units**



High-dimensional input



Large numbers of network parameters

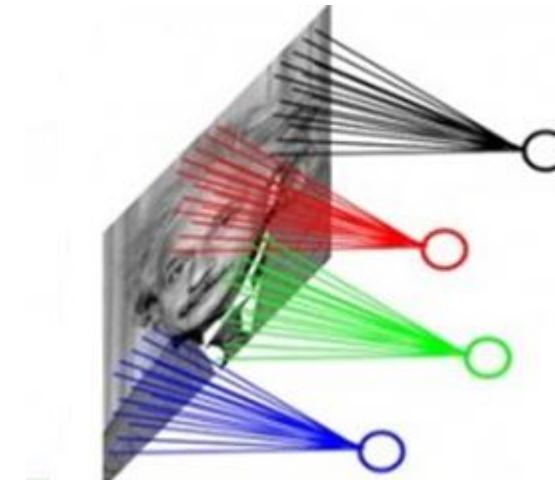
CNN: Local Connection

<http://blog.csdn.net/stdcoutzyx/article/details/41596663>

- Local receptive field, biologically-plausible



Full connection

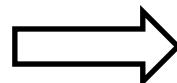


Local connection

- Image filtering



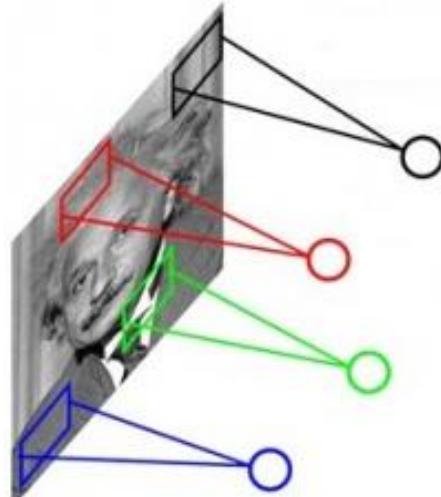
$$\ast \quad \begin{matrix} & & \\ & & \\ \bullet & & \\ & & \\ & & \end{matrix} \quad w_9$$



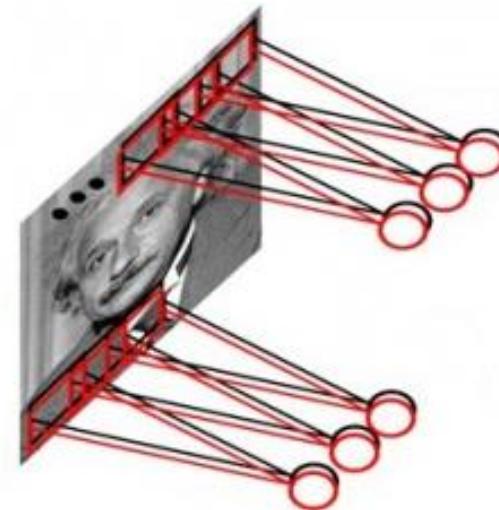
$$y = p_1 w_1 + \cdots + p_9 w_9$$

CNN: Weight Sharing

- All local regions share the same filter weights



Local connection



Weight sharing

- Image convolution

| | | | | |
|---------|---------|---------|---|---|
| 1 x1 | 1 x0 | 1 x1 | 0 | 0 |
| 0 x0 | 1 x1 | 1 x0 | 1 | 0 |
| 0 x1 | 0 x0 | 1 x1 | 1 | 1 |
| 0 | 0 | 1 | 1 | 0 |
| 0 | 1 | 1 | 0 | 0 |

Image

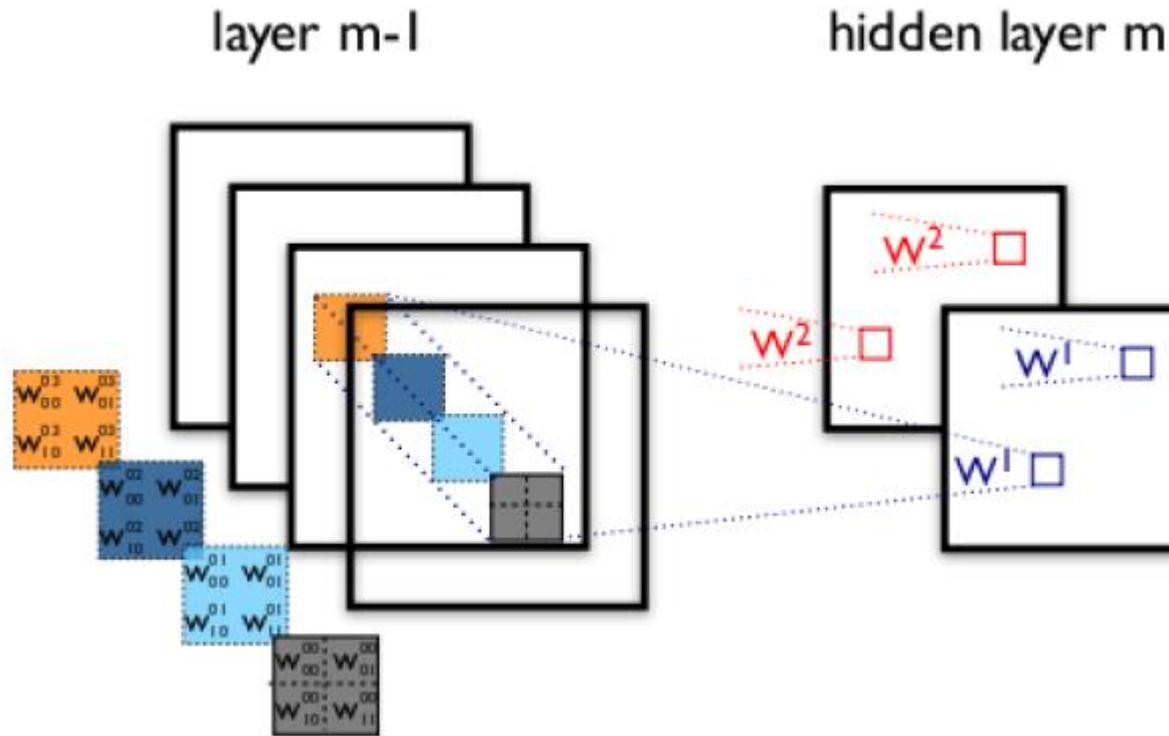
| | | |
|---|--|--|
| 4 | | |
| | | |
| | | |
| | | |

Convolved Feature

$$h_{ij}^k = \tanh((W^k * x)_{ij} + b_k)$$

CNN: Multiple Filters

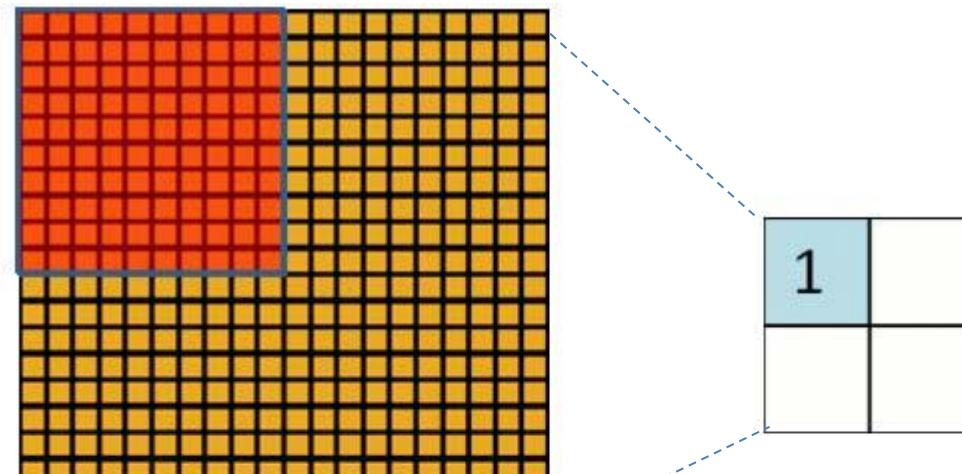
- Use N different filters obtain N feature maps



w_1 : weights of the 1st filter, w_2 : weights of the 2nd filter

CNN: Pooling

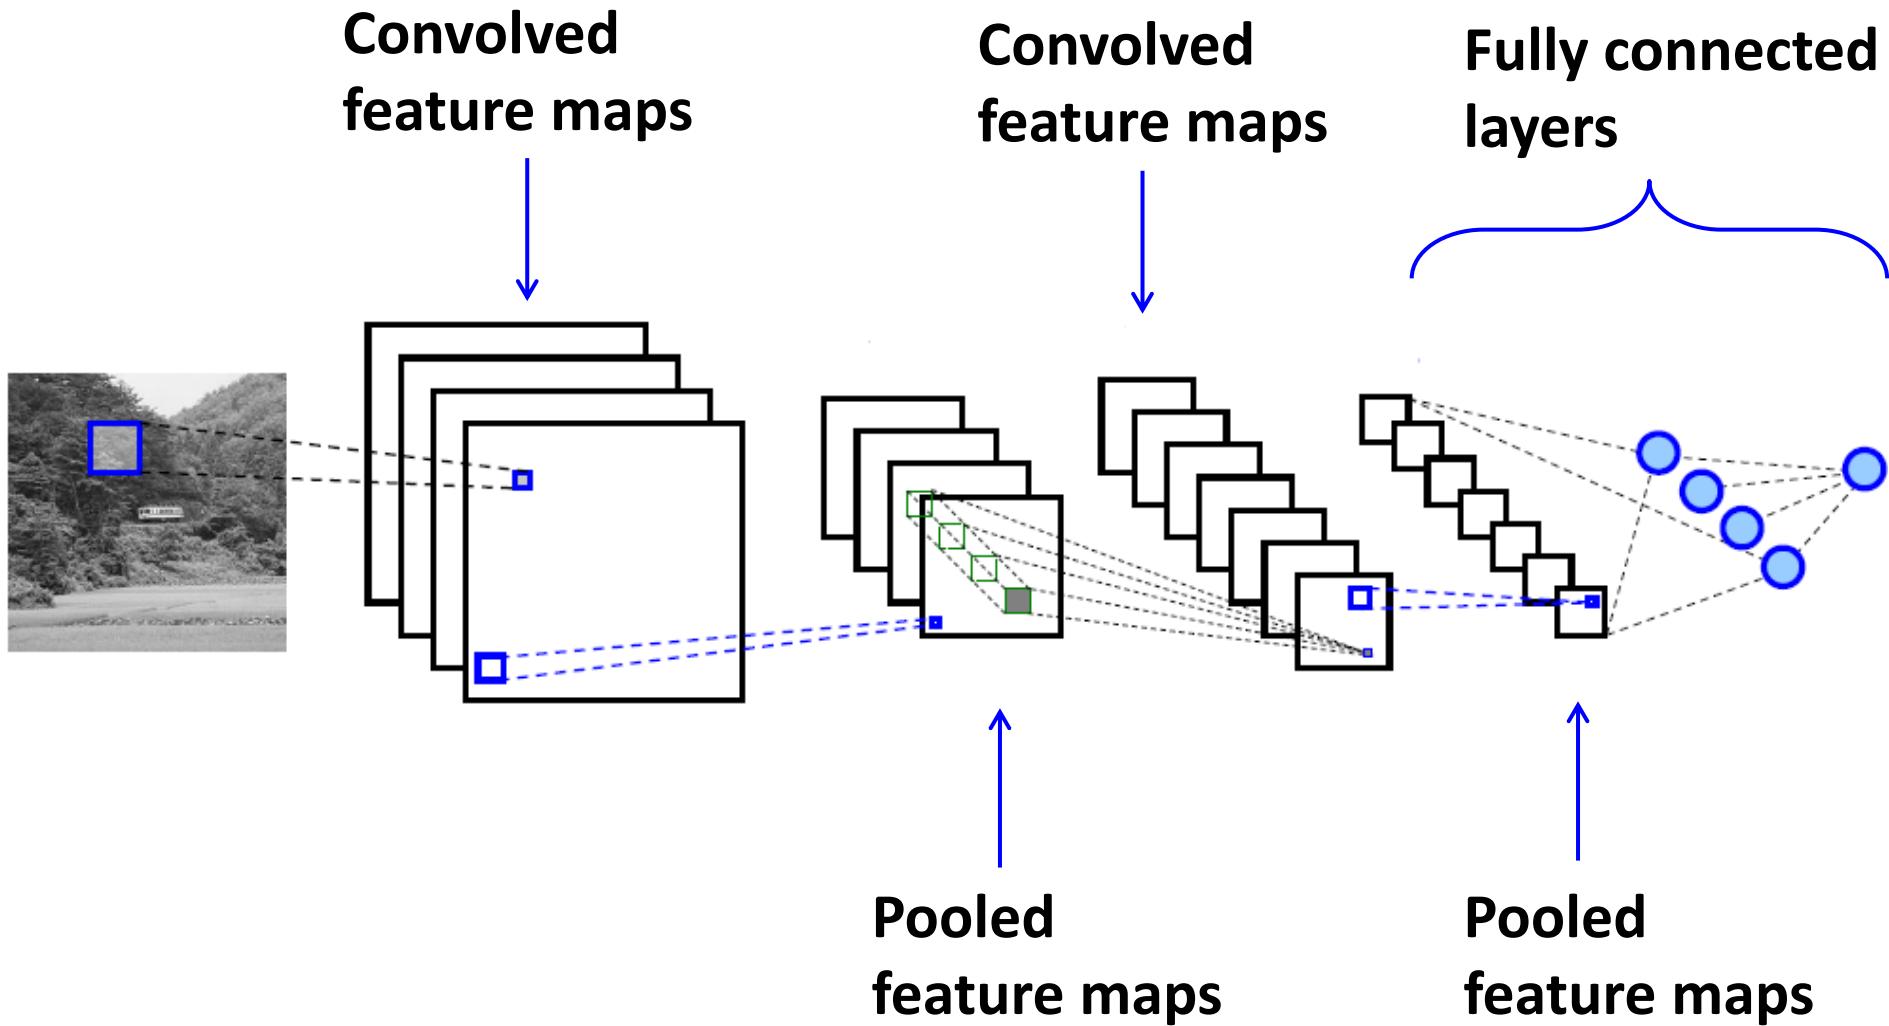
- Be invariant to image shifting using max (or average) pooling



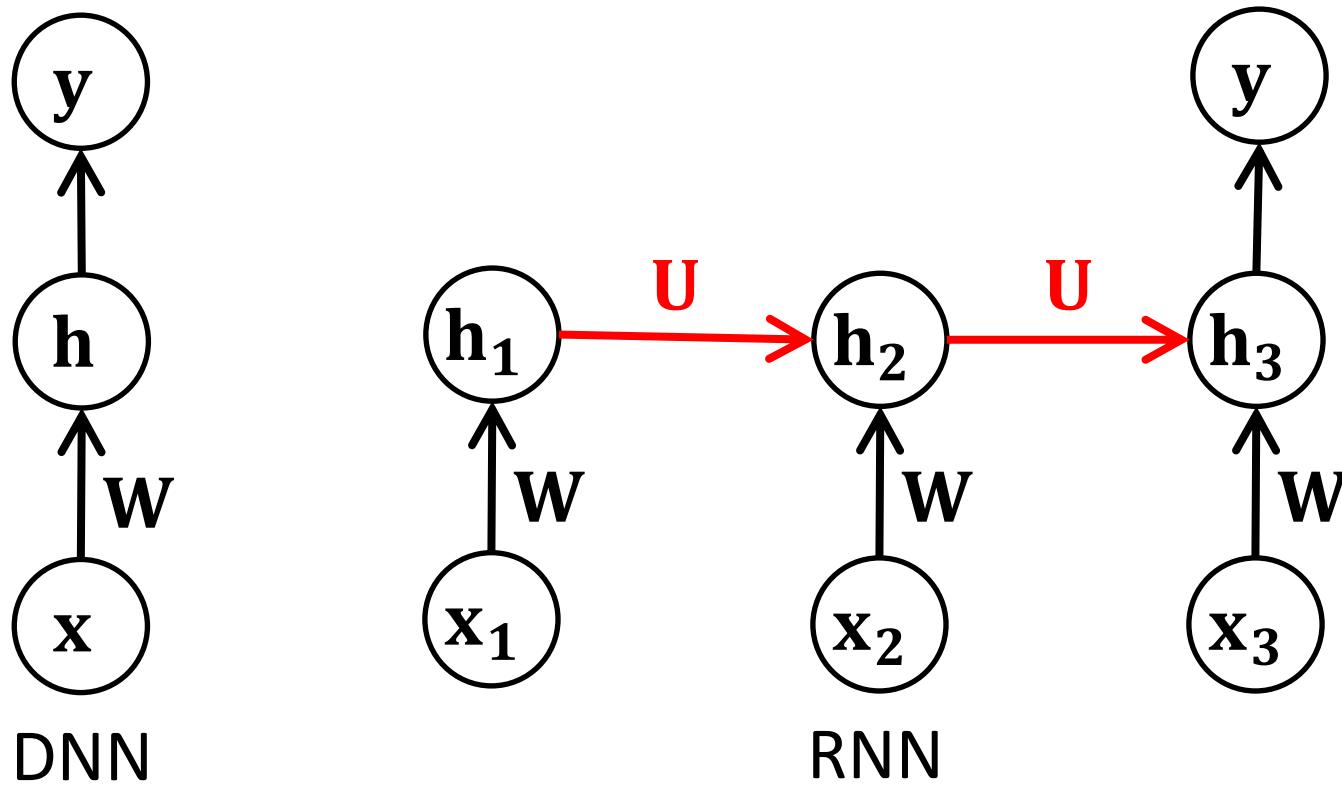
Convolved feature

Pooled feature

The Whole Architecture

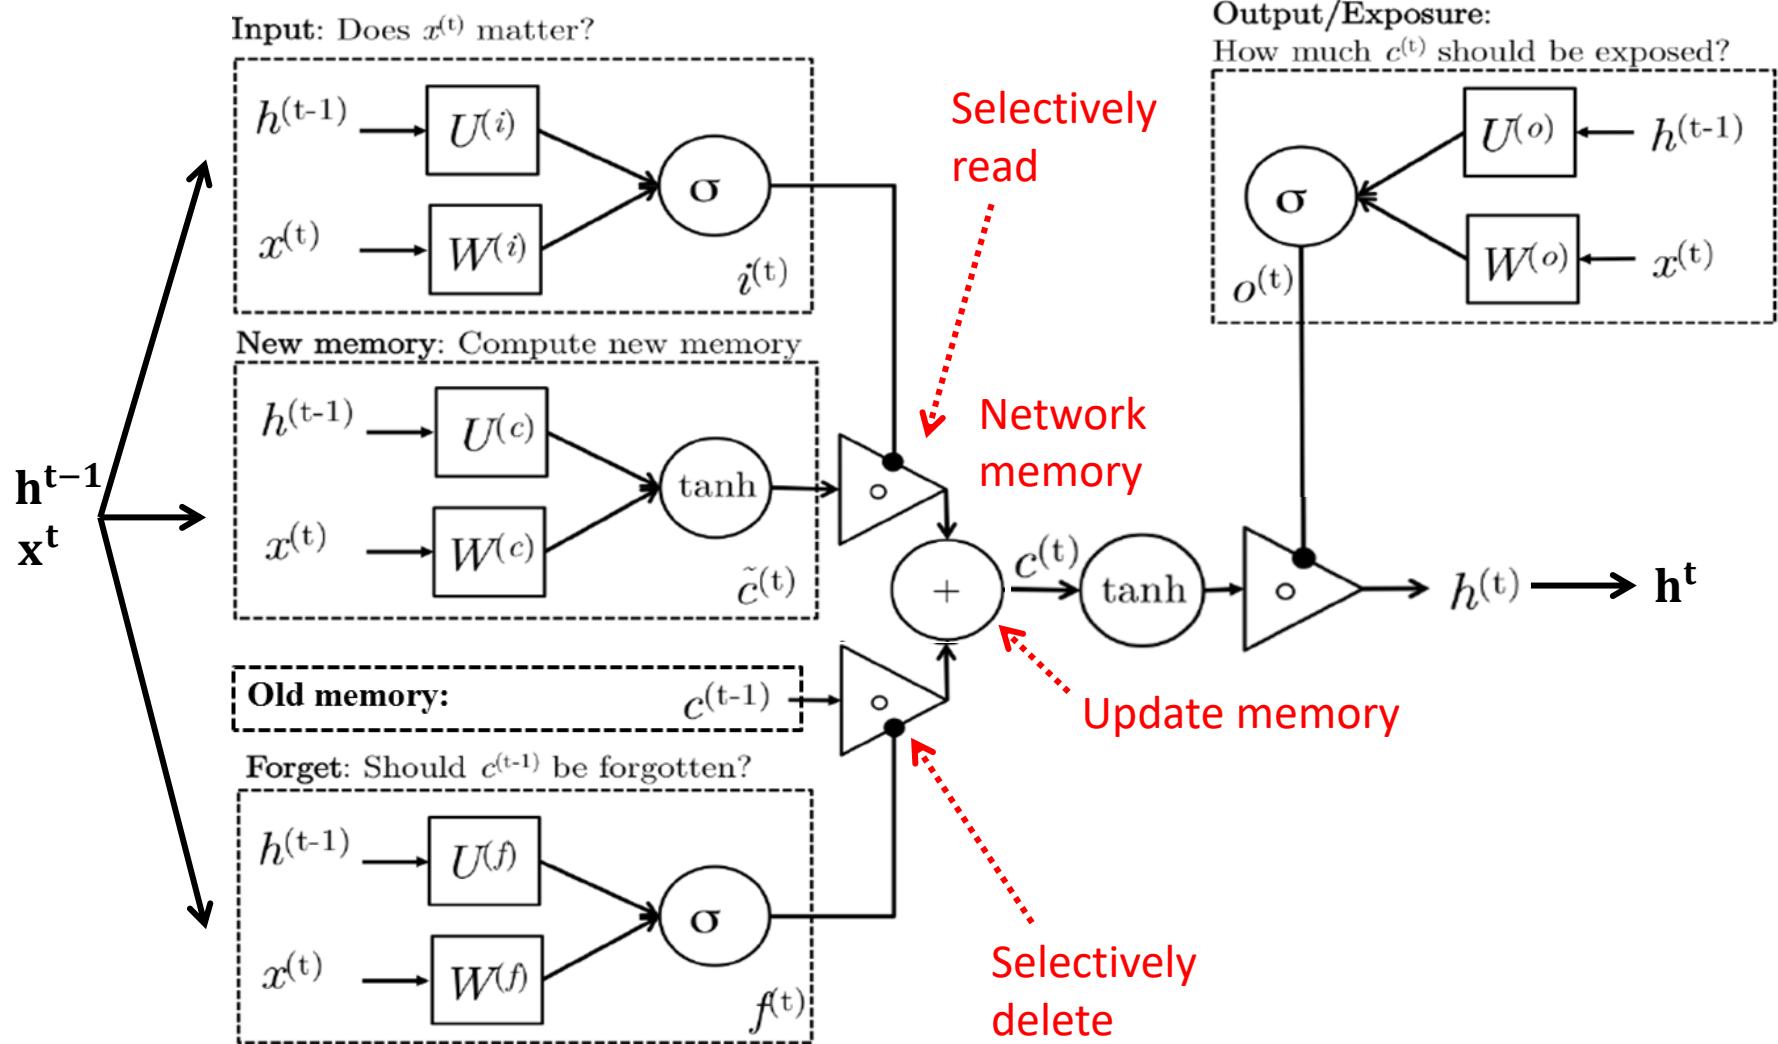


RNN



- $\mathbf{x}_t \in \mathbb{R}^d, \mathbf{h}_t \in \mathbb{R}^n, \mathbf{W} \in \mathbb{R}^{d \times n}, \mathbf{U} \in \mathbb{R}^{n \times n}$
- \mathbf{h}_t RNN aim to model long-range contextual information in forward direction

LSTM



[Hochreiter and Schmidhuber, Neural computation 1997]

Outline

1 Course Information

2 What is Deep Learning

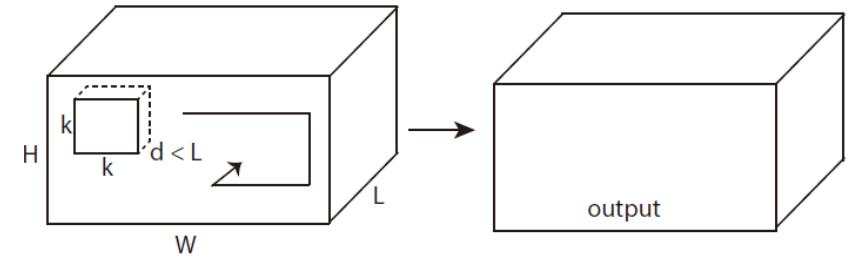
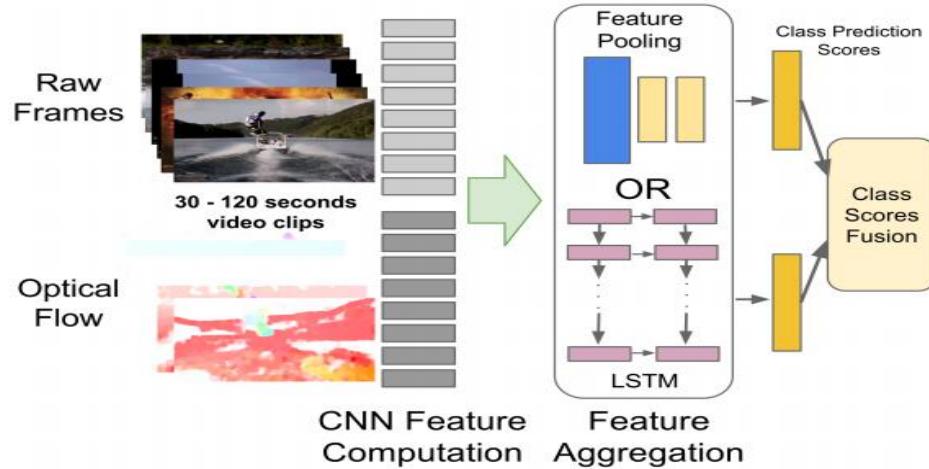
3 Applications of Deep Learning

4 Future Directions

5 Project and homework guidance

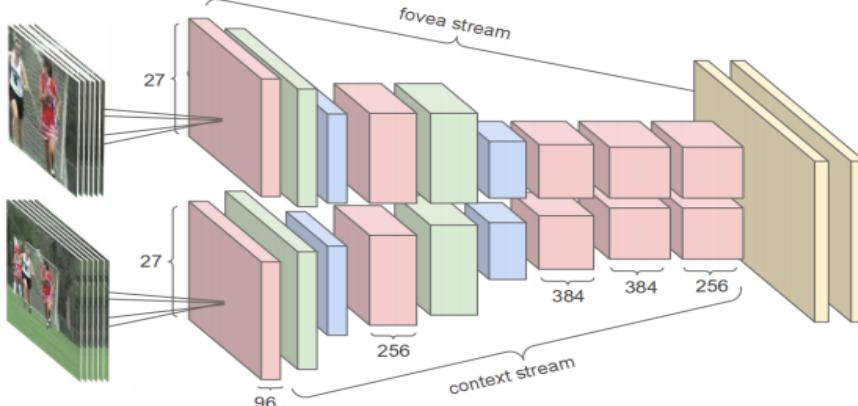
Application for Video Classification

CNN: convolutional neural networks; RNN: recurrent neural networks

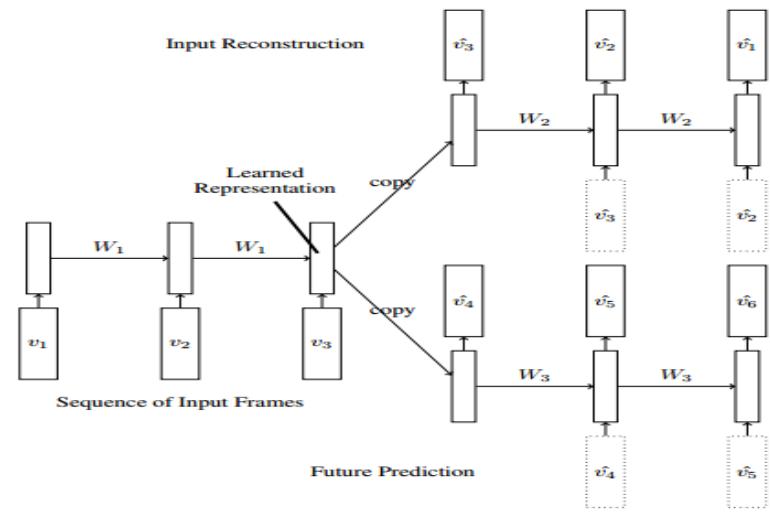


3D Convolution, ICCV2015

Long-range Videos Modelling, CVPR2015



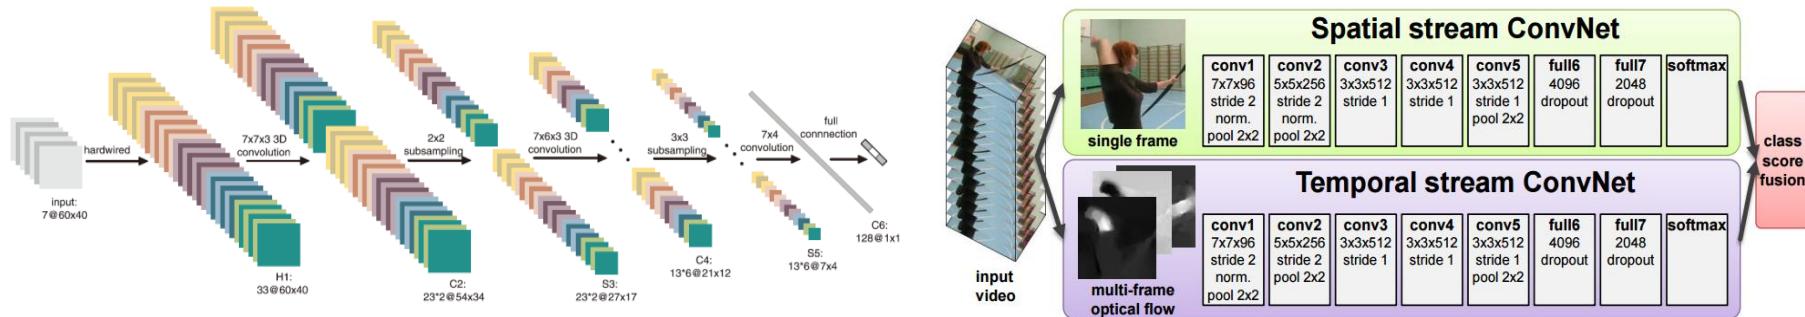
Multi-resolution CNN, CVPR2014



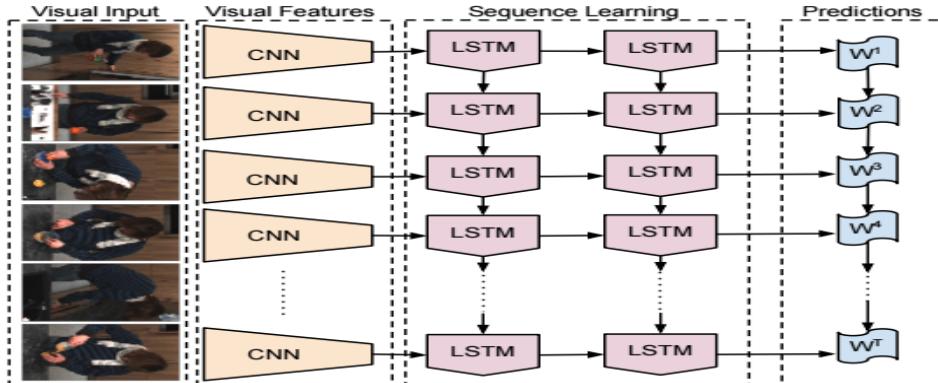
Autoencoder-RNN, ICML2015

Application for Action Recognition

CNN: convolutional neural networks; LSTM: long short term memory; RNN: recurrent neural networks

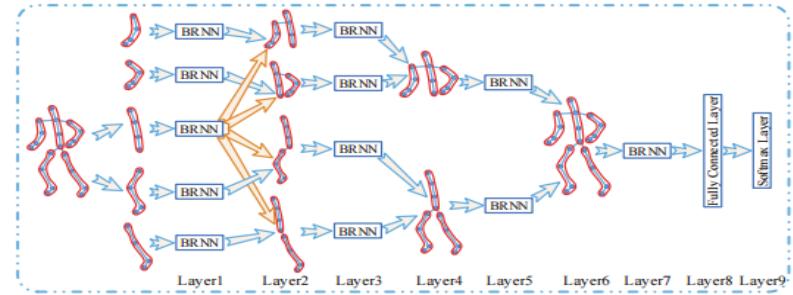


3D CNN, ICML2010



CNN + LSTM-RNN, CVPR2015

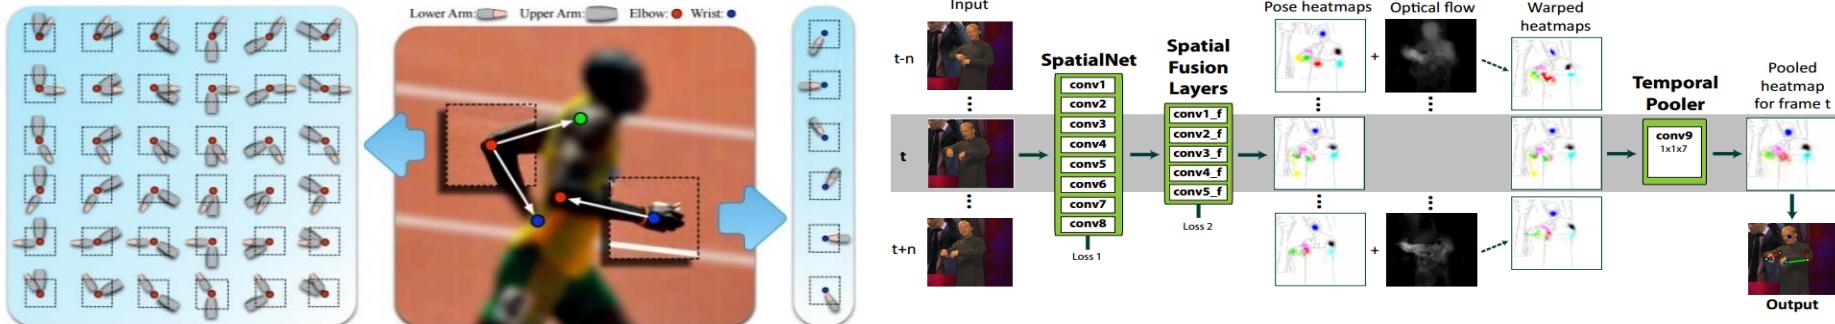
Two-stream CNN, NeurIPS2014



Skeleton + LSTM-RNN, CVPR2015

Application for Pose Estimation

CNN: convolutional neural networks; DNN: deep neural networks



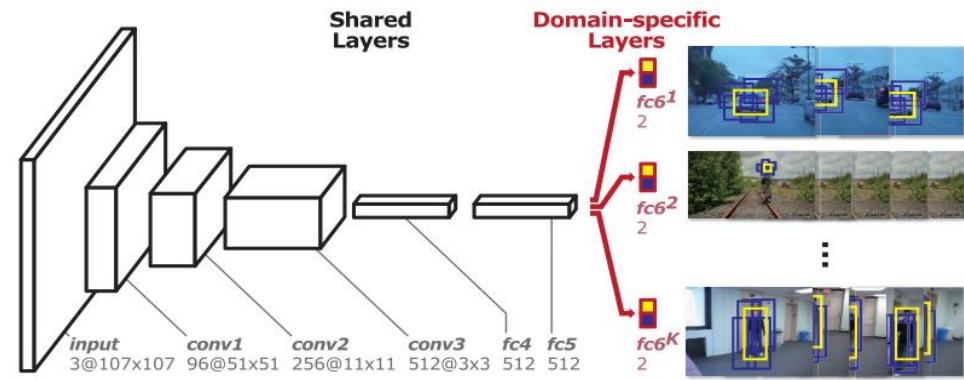
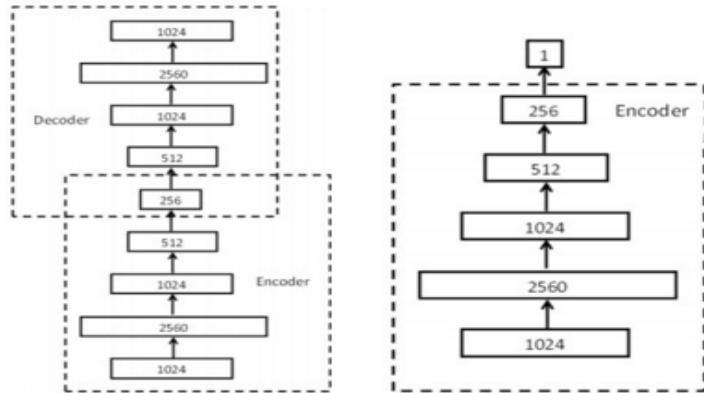
CNN + Graphical Model, NeurIPS2014

Flowing ConvNets, ICCV2015



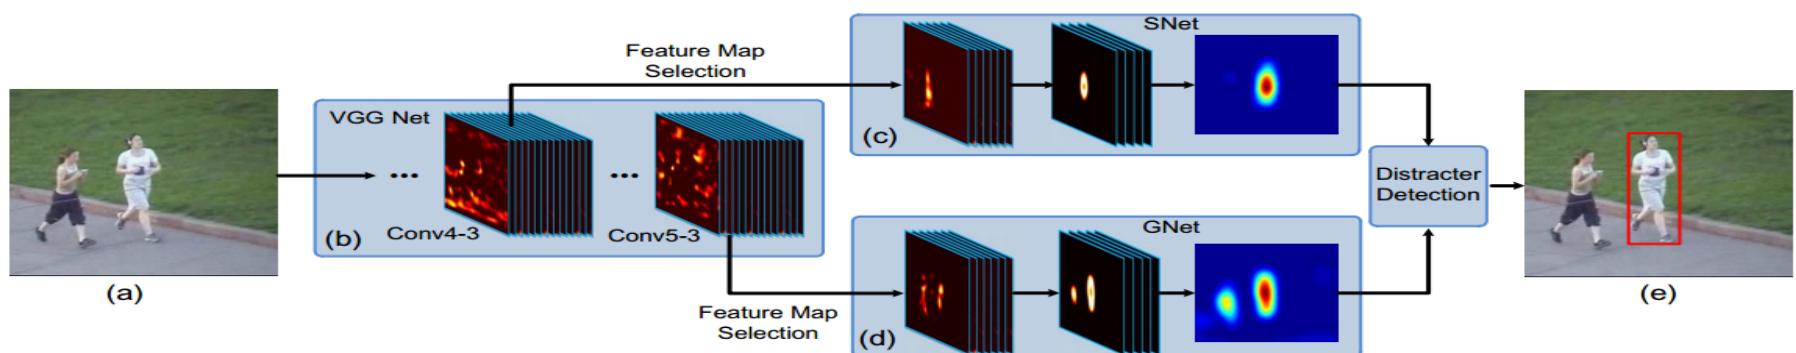
DNN-based Regression, CVPR2014

Application for Video Tracking



Stacked Denoising Auto-encoder, NeurIPS2013

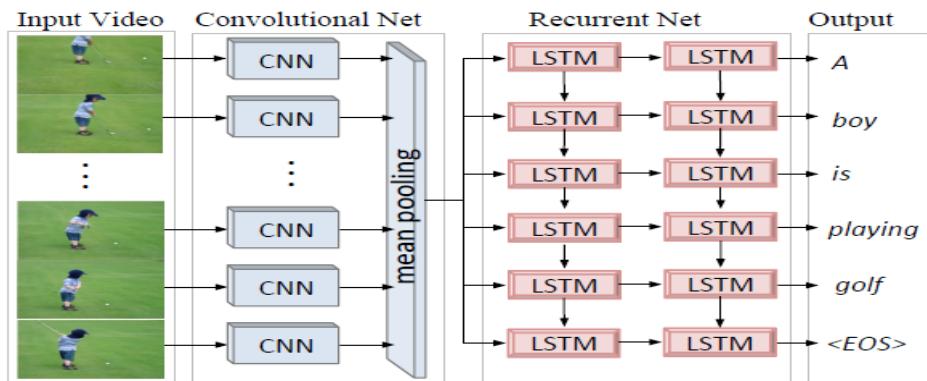
Multi-domain Networks, CVPR2016



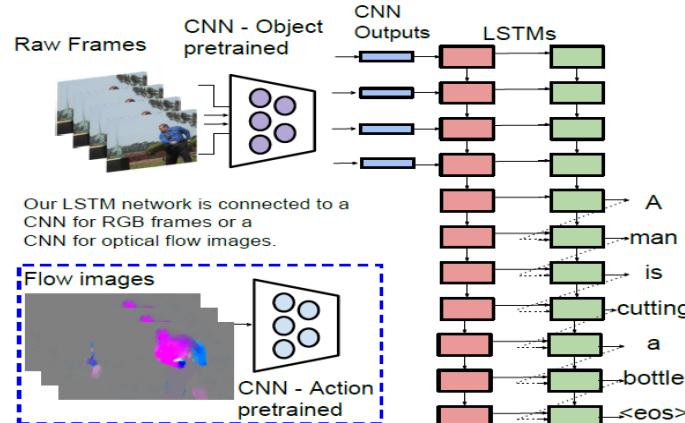
Fully Convolutional Networks, ICCV2015

Application for Video Caption

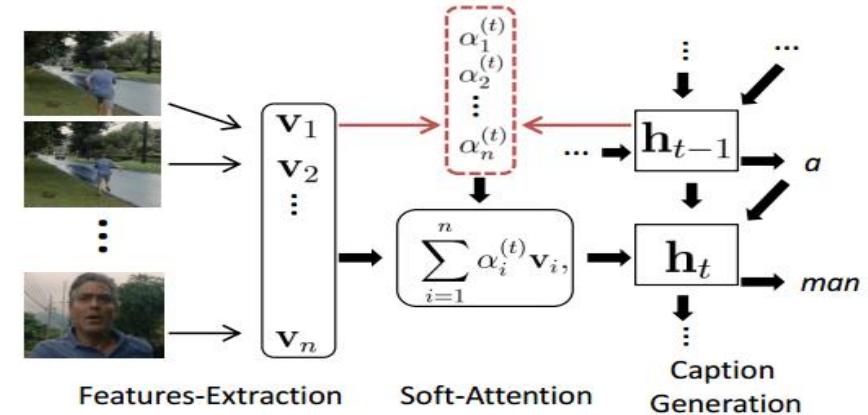
CNN: convolutional neural networks; LSTM: long short term memory; RNN: recurrent neural networks



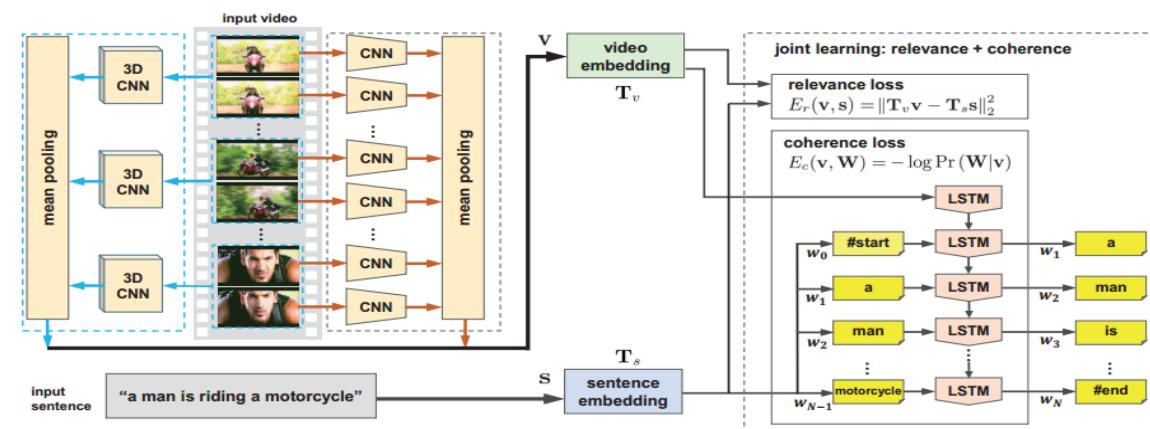
CNN + LSTM-RNN, CVPR2015



Video2Text, ICCV2015

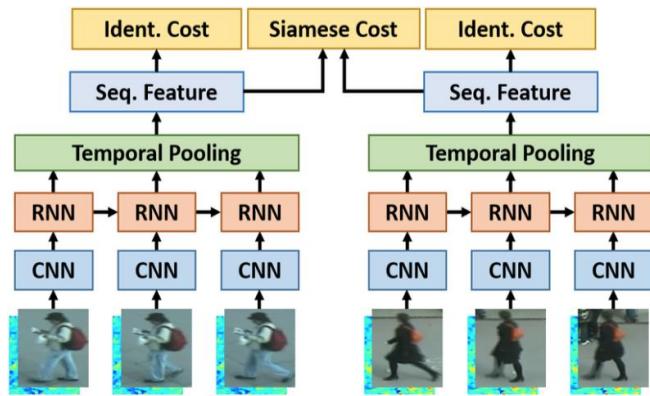


Attention-based LSTM-RNN, ICCV2015

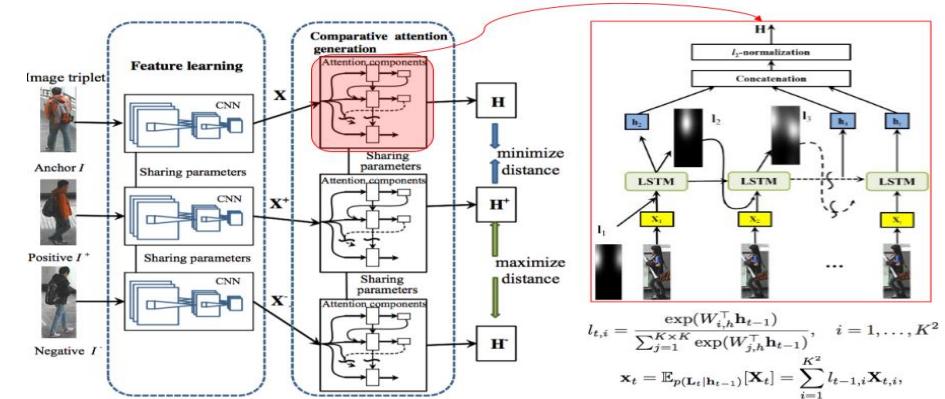


Visual-semantic Embedding + LSTM-RNN, CVPR2016

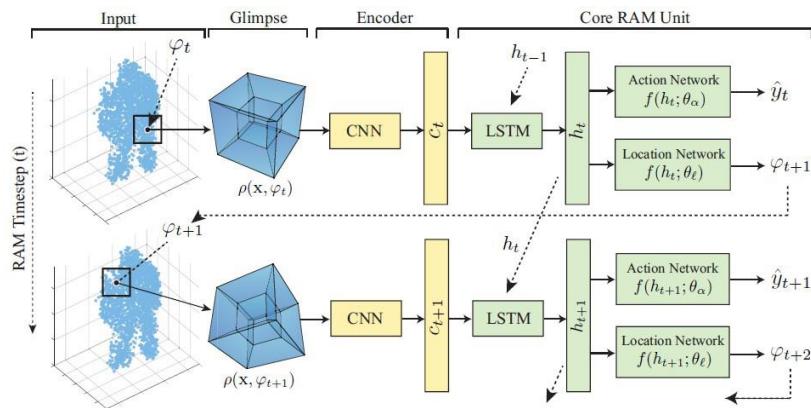
Application for Person Re-identification



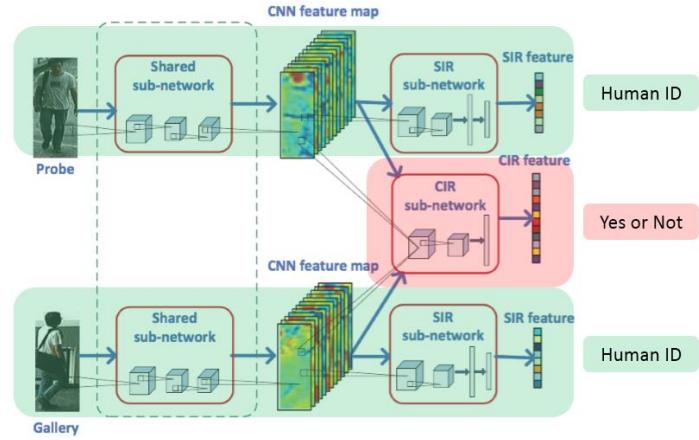
Recurrent CNN, CVPR2016



End-to-end Attention, TIP2016

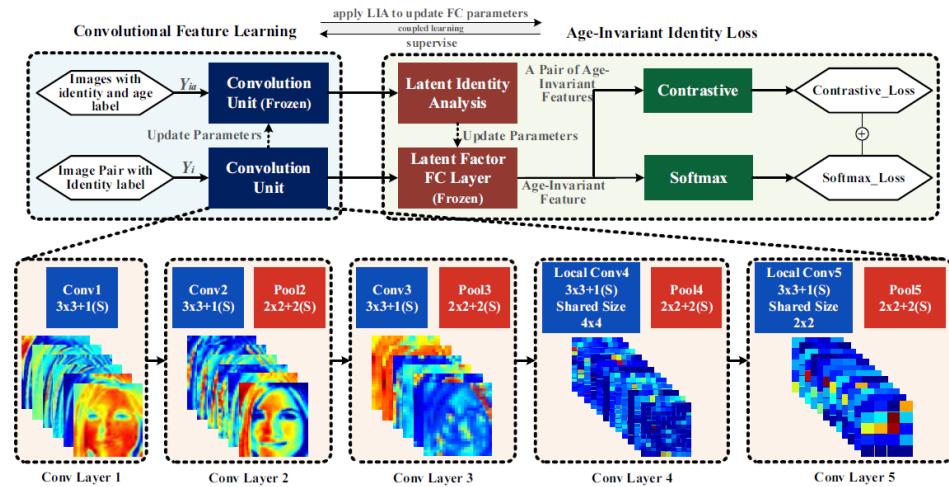
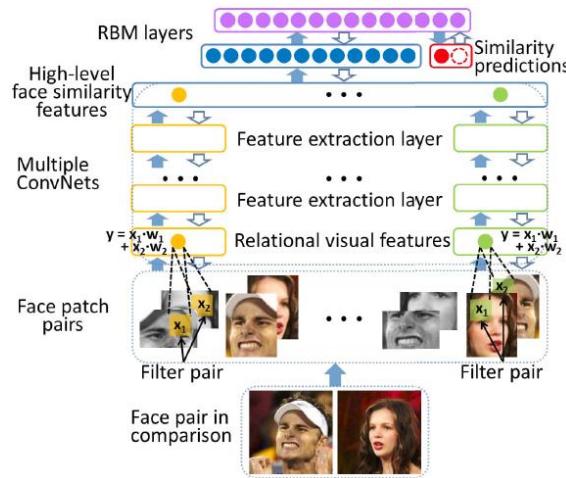


Sequential Local Attention, CVPR2016

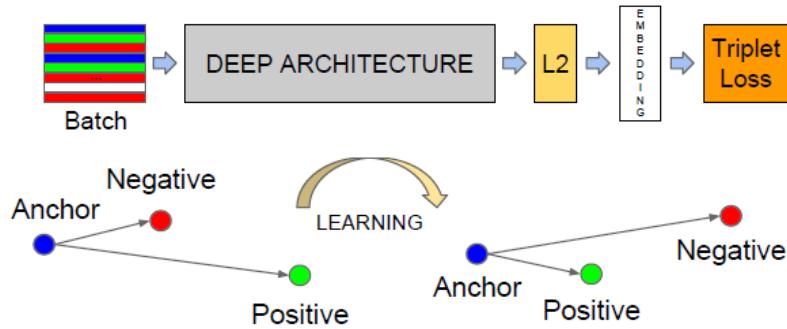


Single-/Cross-image Representations, CVPR2016

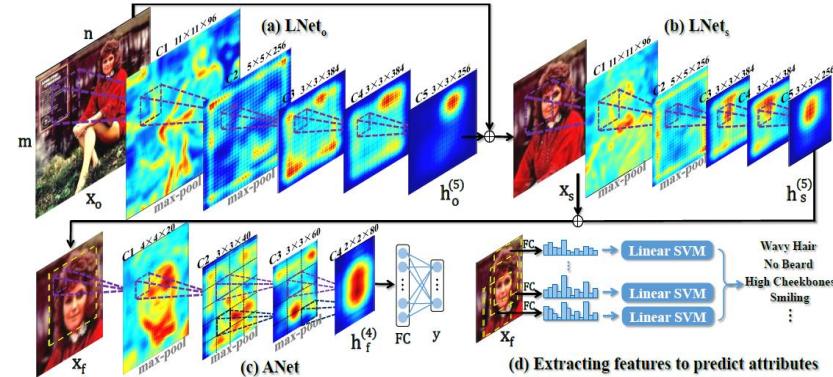
Application for Face Recognition



Hybrid ConvNet-RBM, ICCV2013

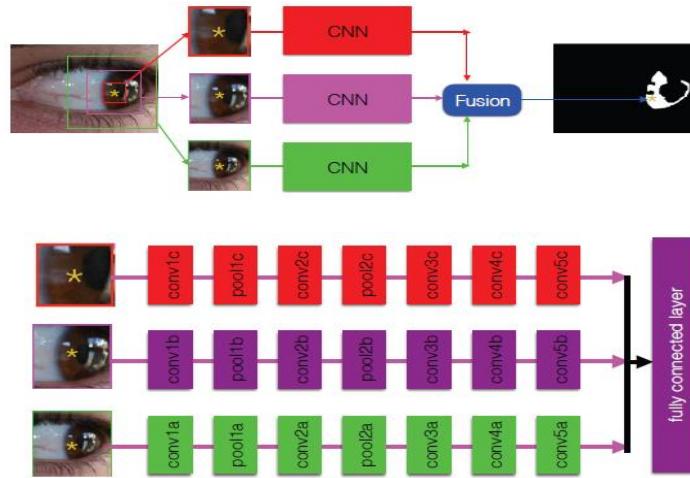


FaceNet, CVPR2015

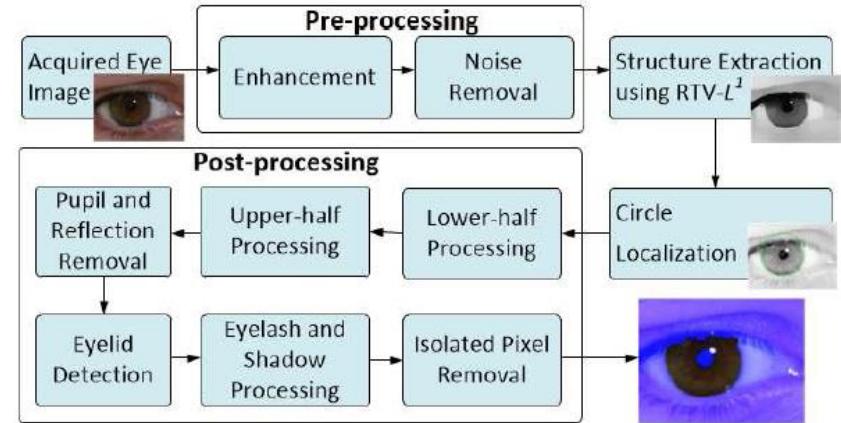


Face LNet and ANet, ICCV2015

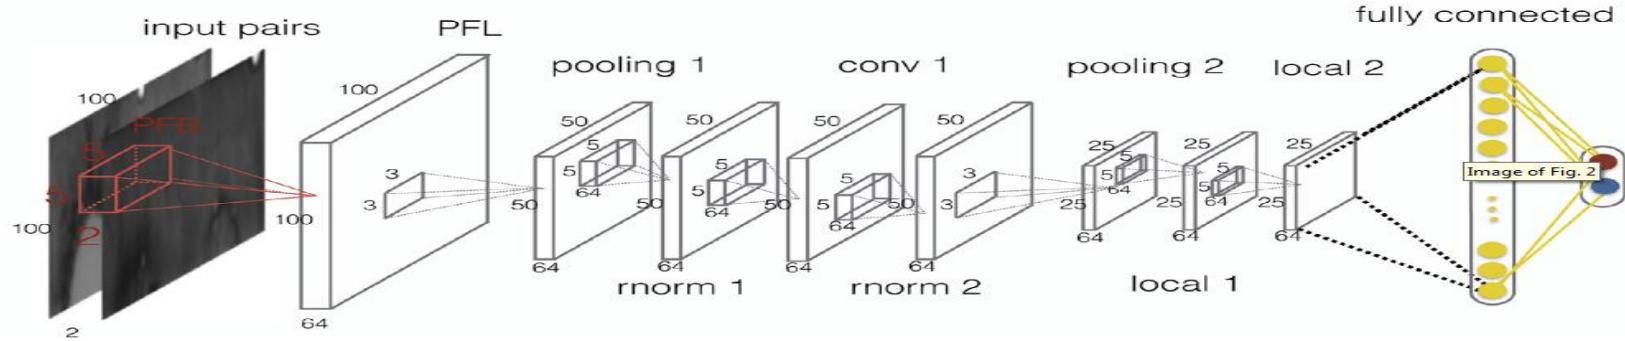
Application for Iris Recognition



HCNNs, ICB2016

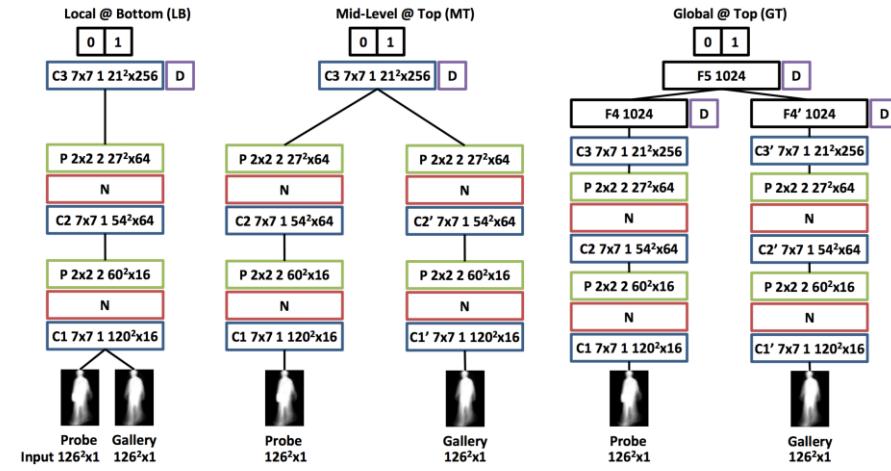
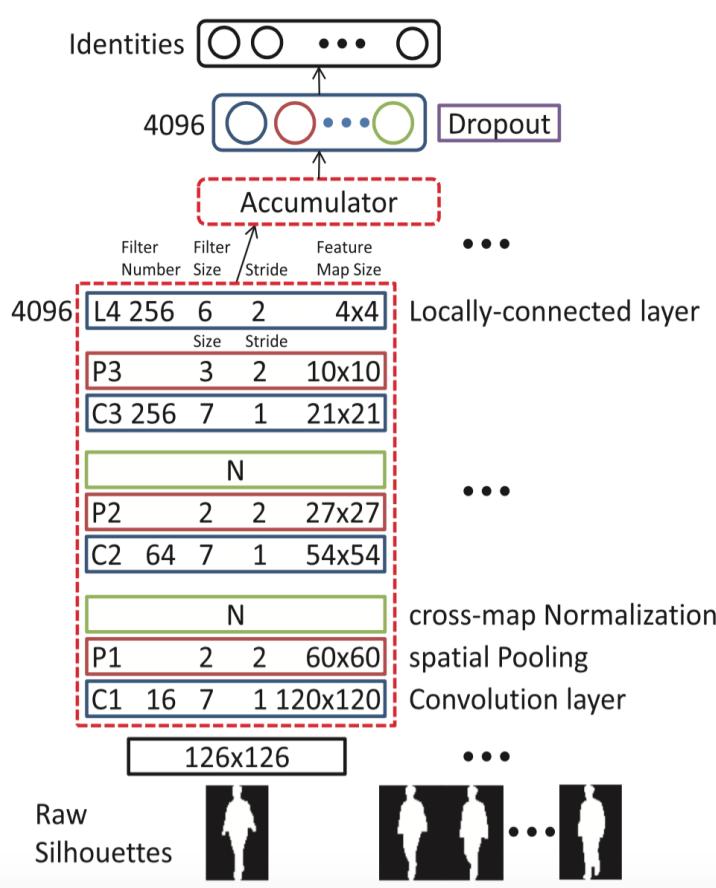


Iris Segmentation, ICCV2015

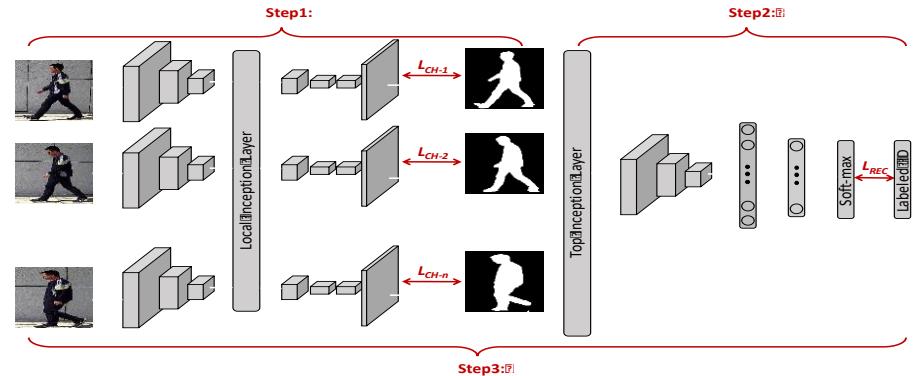


Deepliris, PRL2015

Application for Gait Recognition



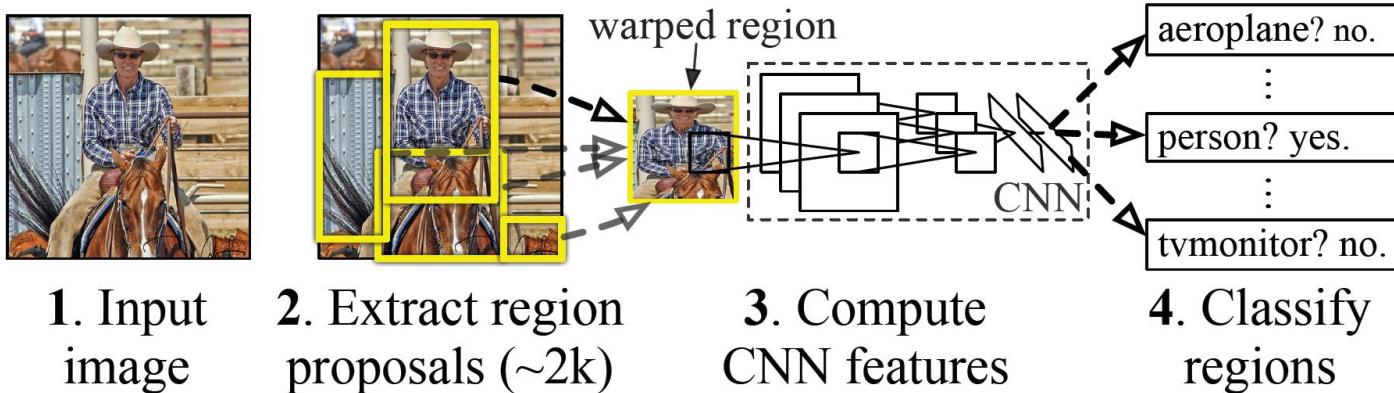
Cross-view Gait Identification, TPAMI 2016



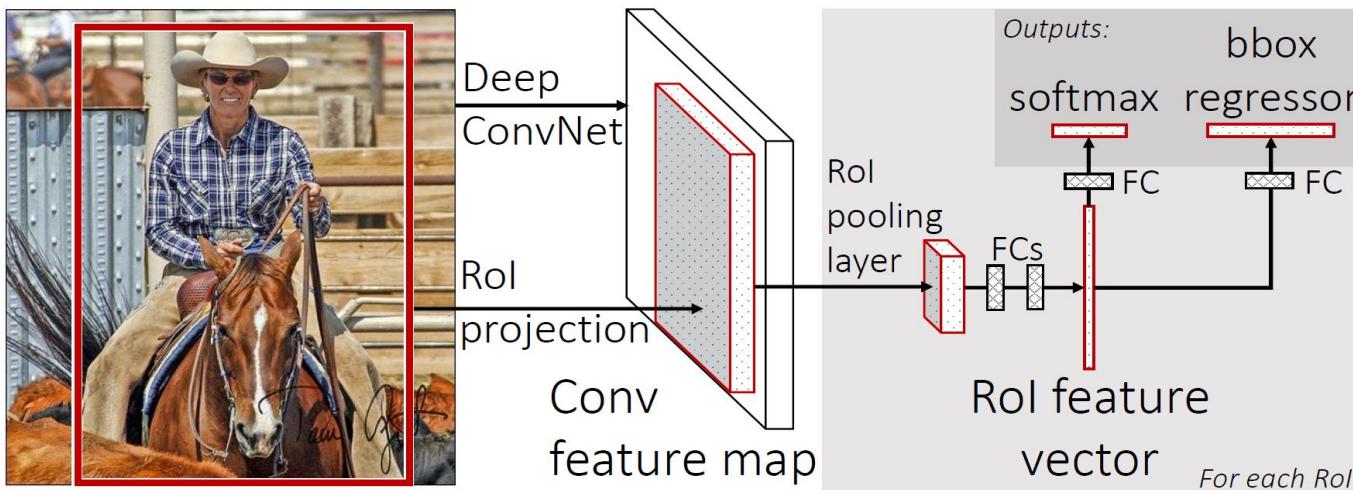
Deep Gait Feature Learning by Image Set, submitted to TIP 2017

Joint Gait Segmentation and Recognition, submitted to CVPR 2017

Application for Object Detection

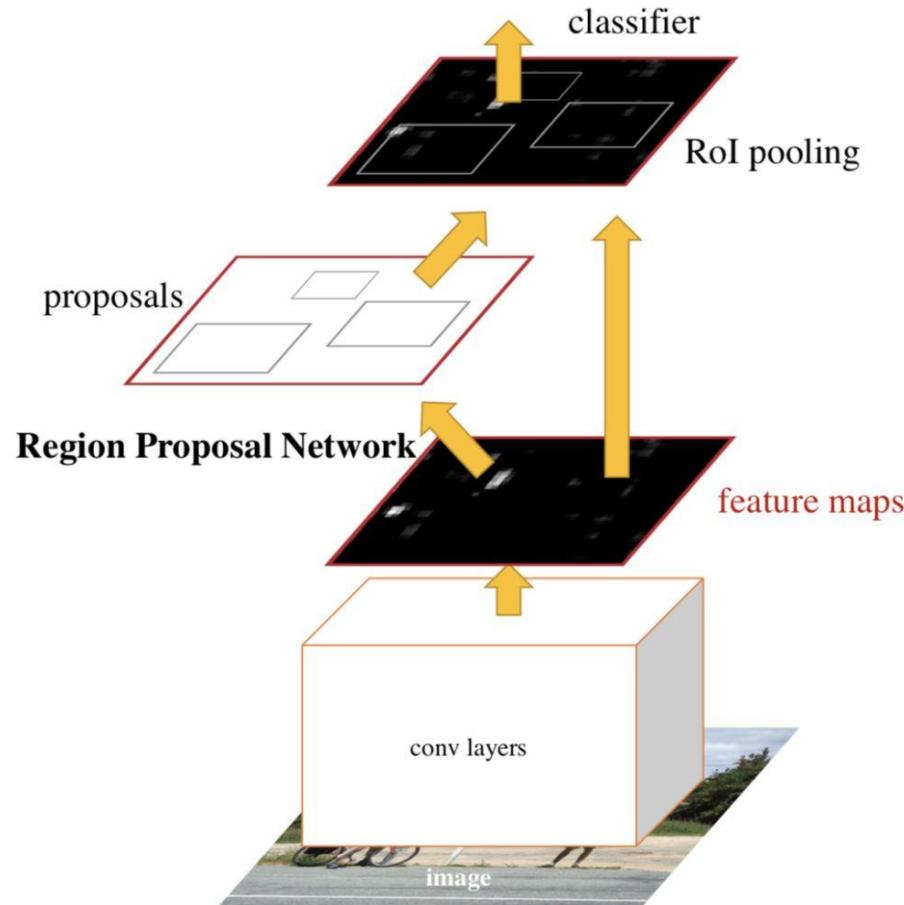


RCNN, CVPR2014



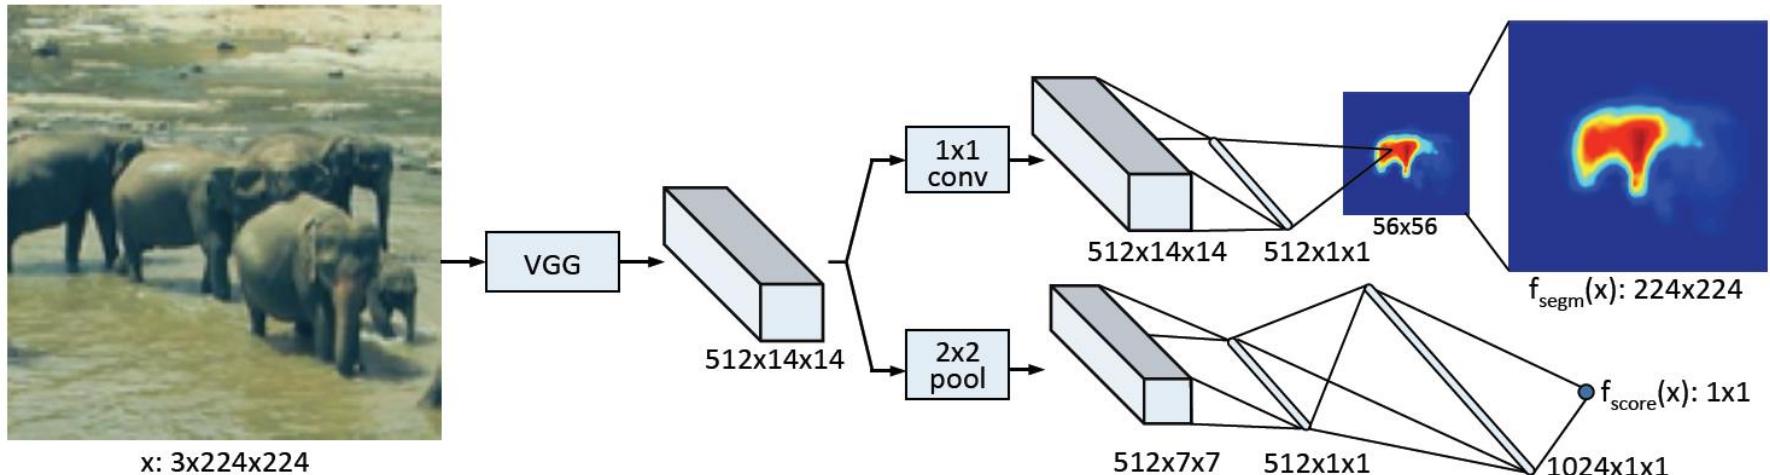
Fast R-CNN, ICCV2015

Application for Object Detection

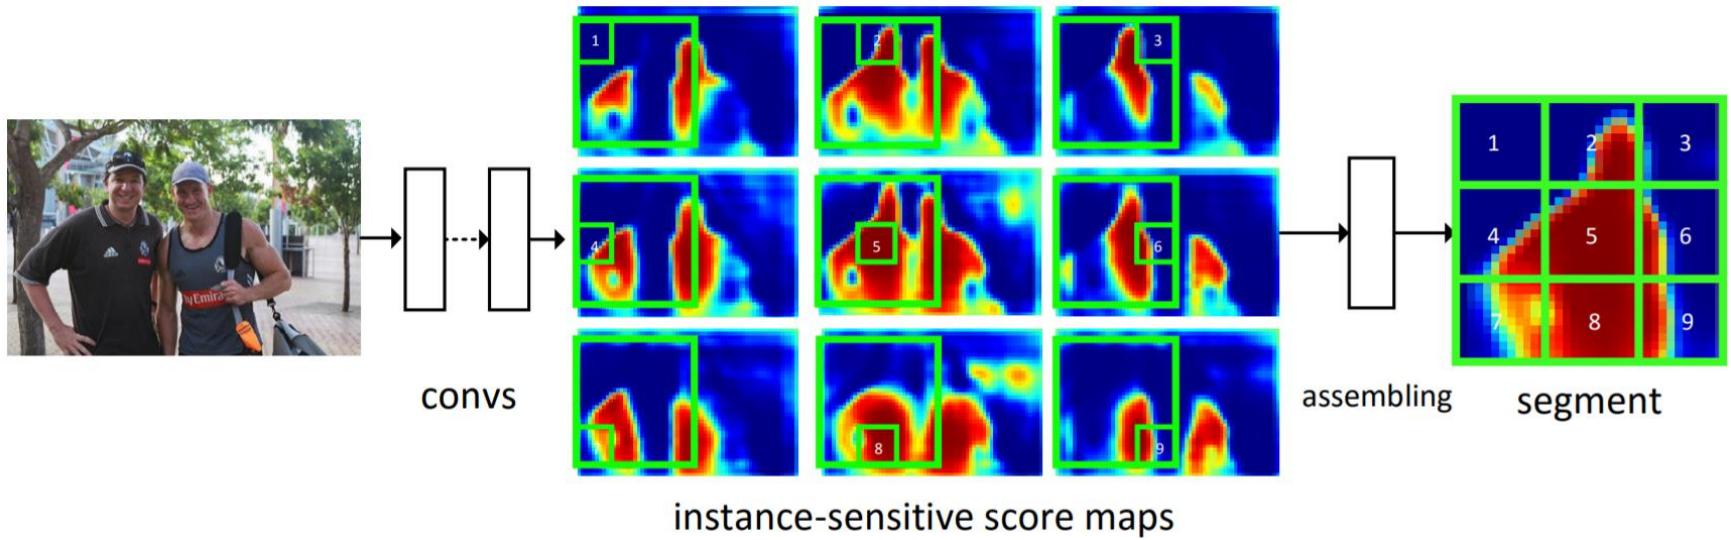


Faster R-CNN, NeurIPS 2015

Application for Semantic Segmentation

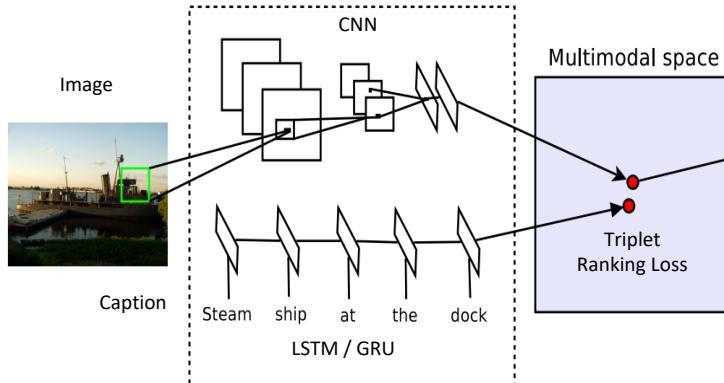


Instance-aware semantic segmentation, CVPR2016

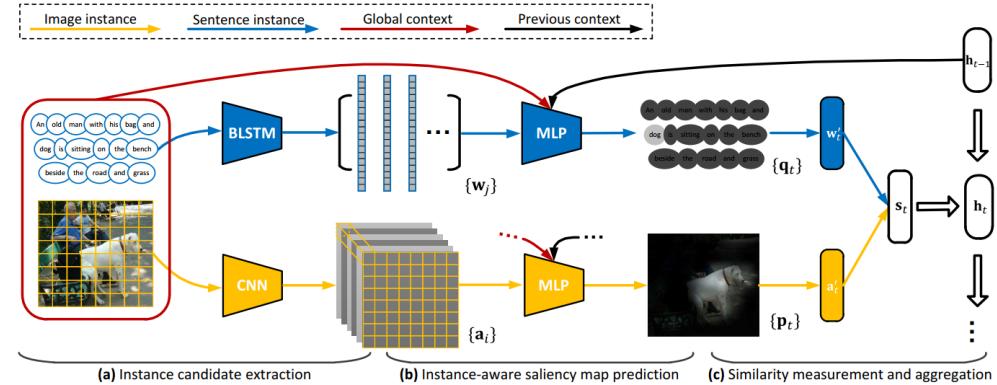


Learning to segment object candidates, NeurIPS2015

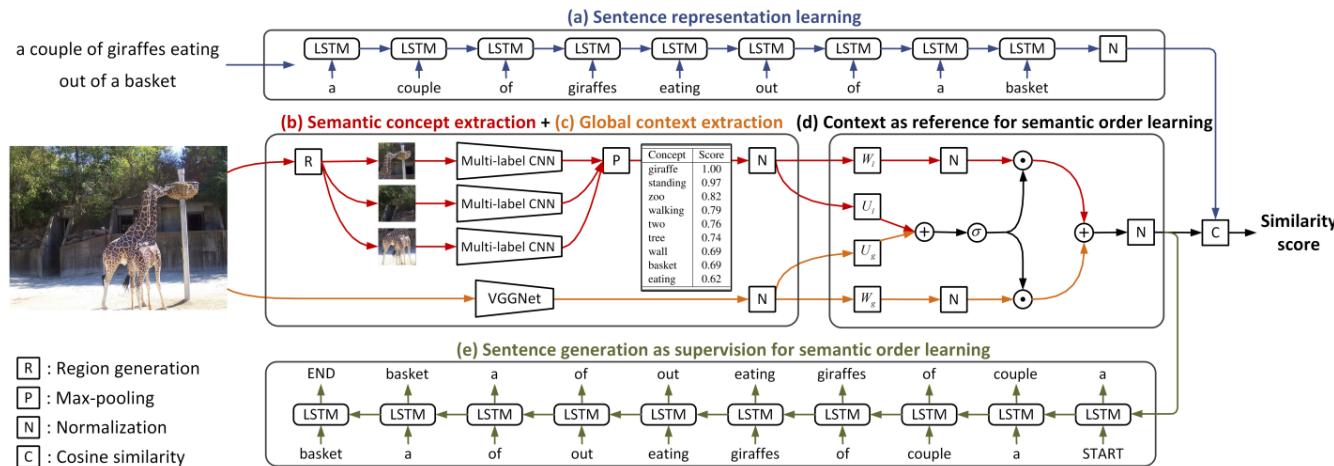
Application for Cross-modal Retrieval



Joint embedding learning, NeurIPS2013

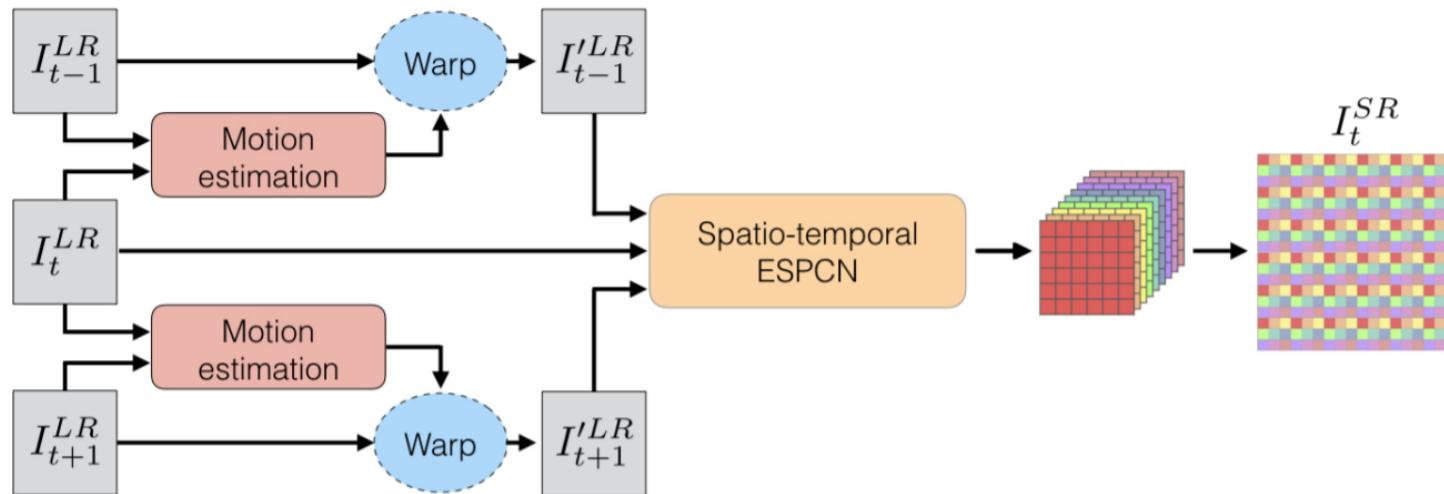


Pair-wise attention alignment, CVPR2017

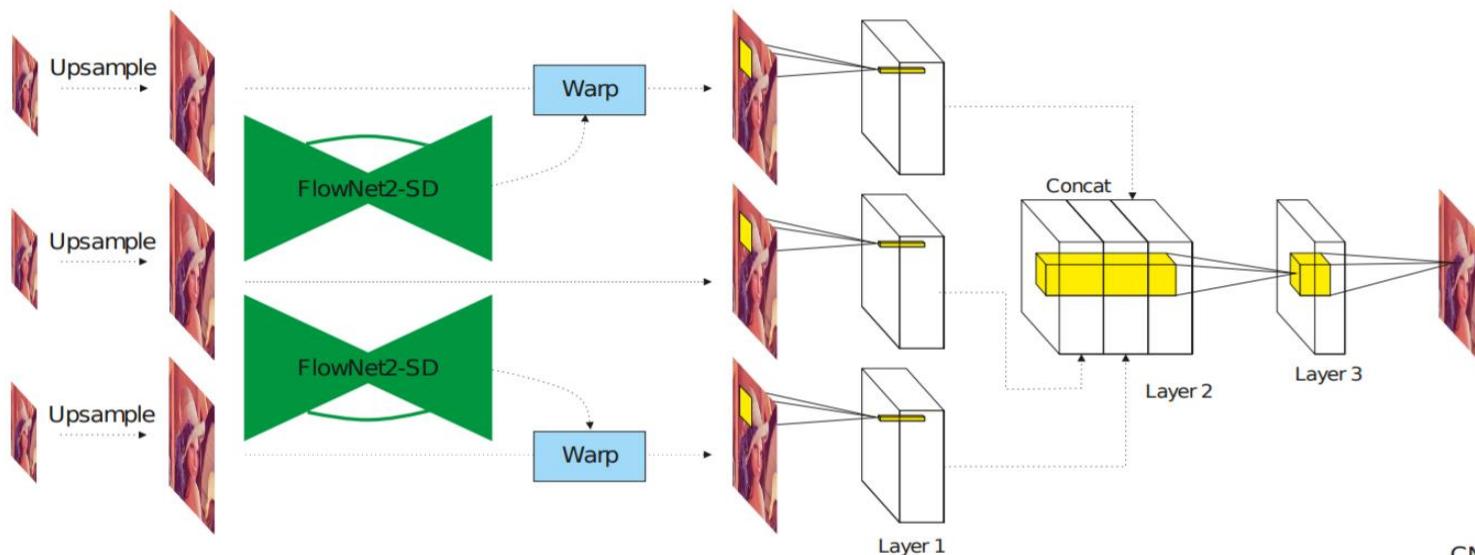


Semantic concept and order learning, CVPR2018

Application for Video Super-resolution



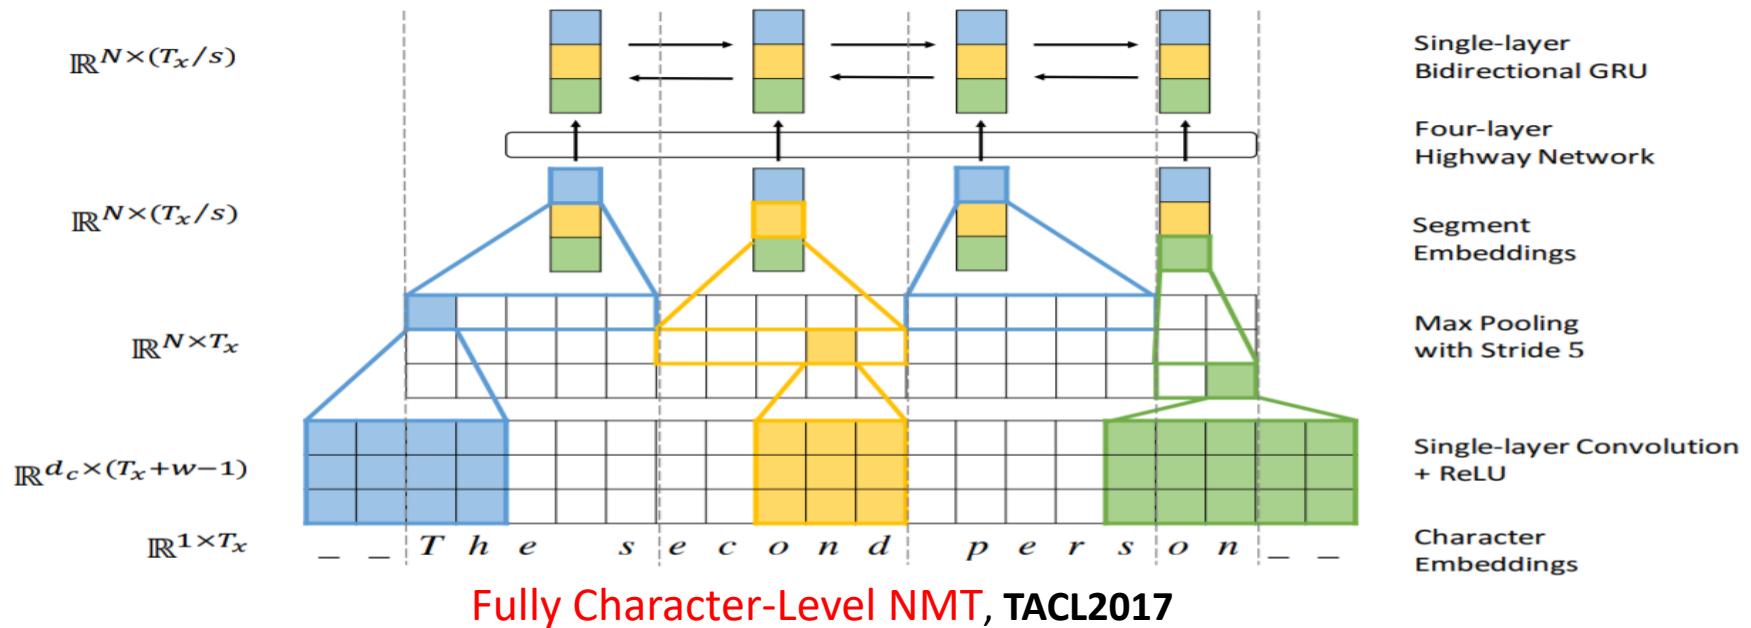
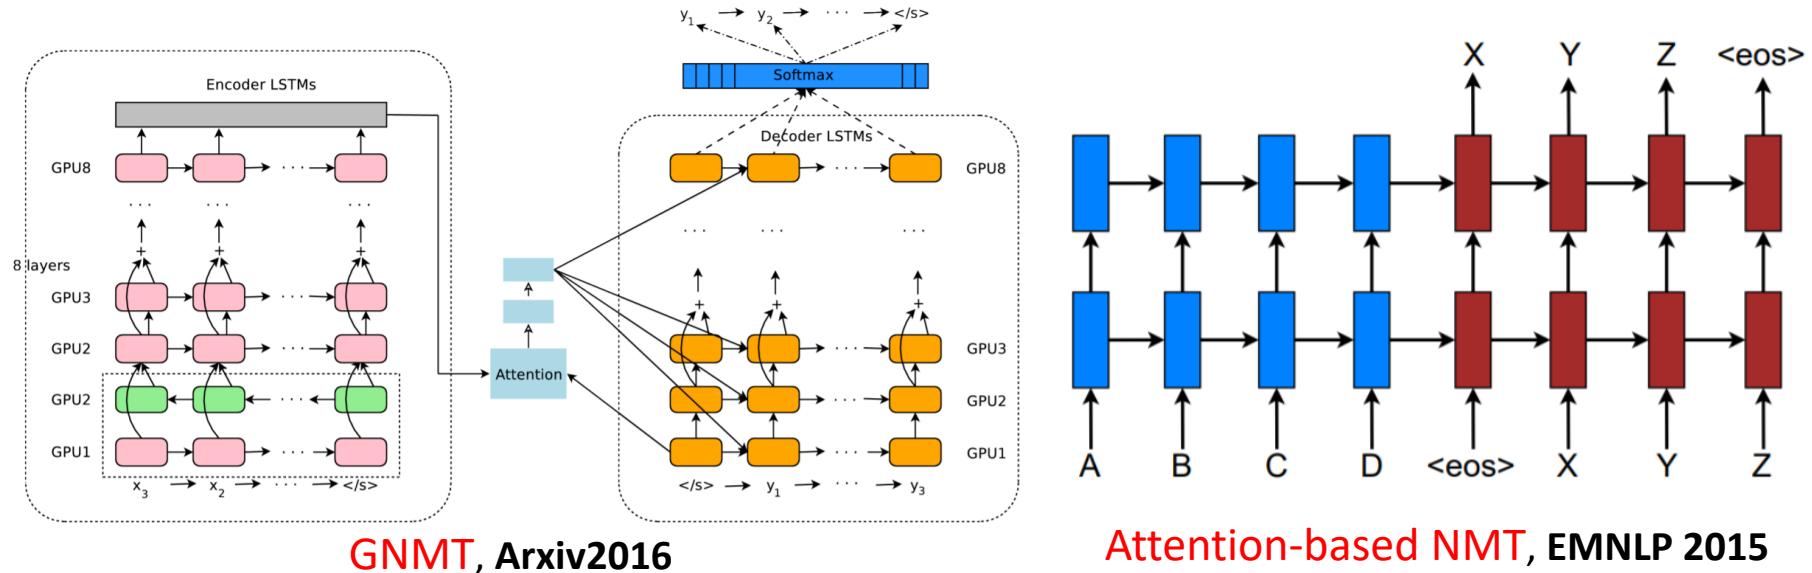
Real-Time Video Super-Resolution, CVPR2017



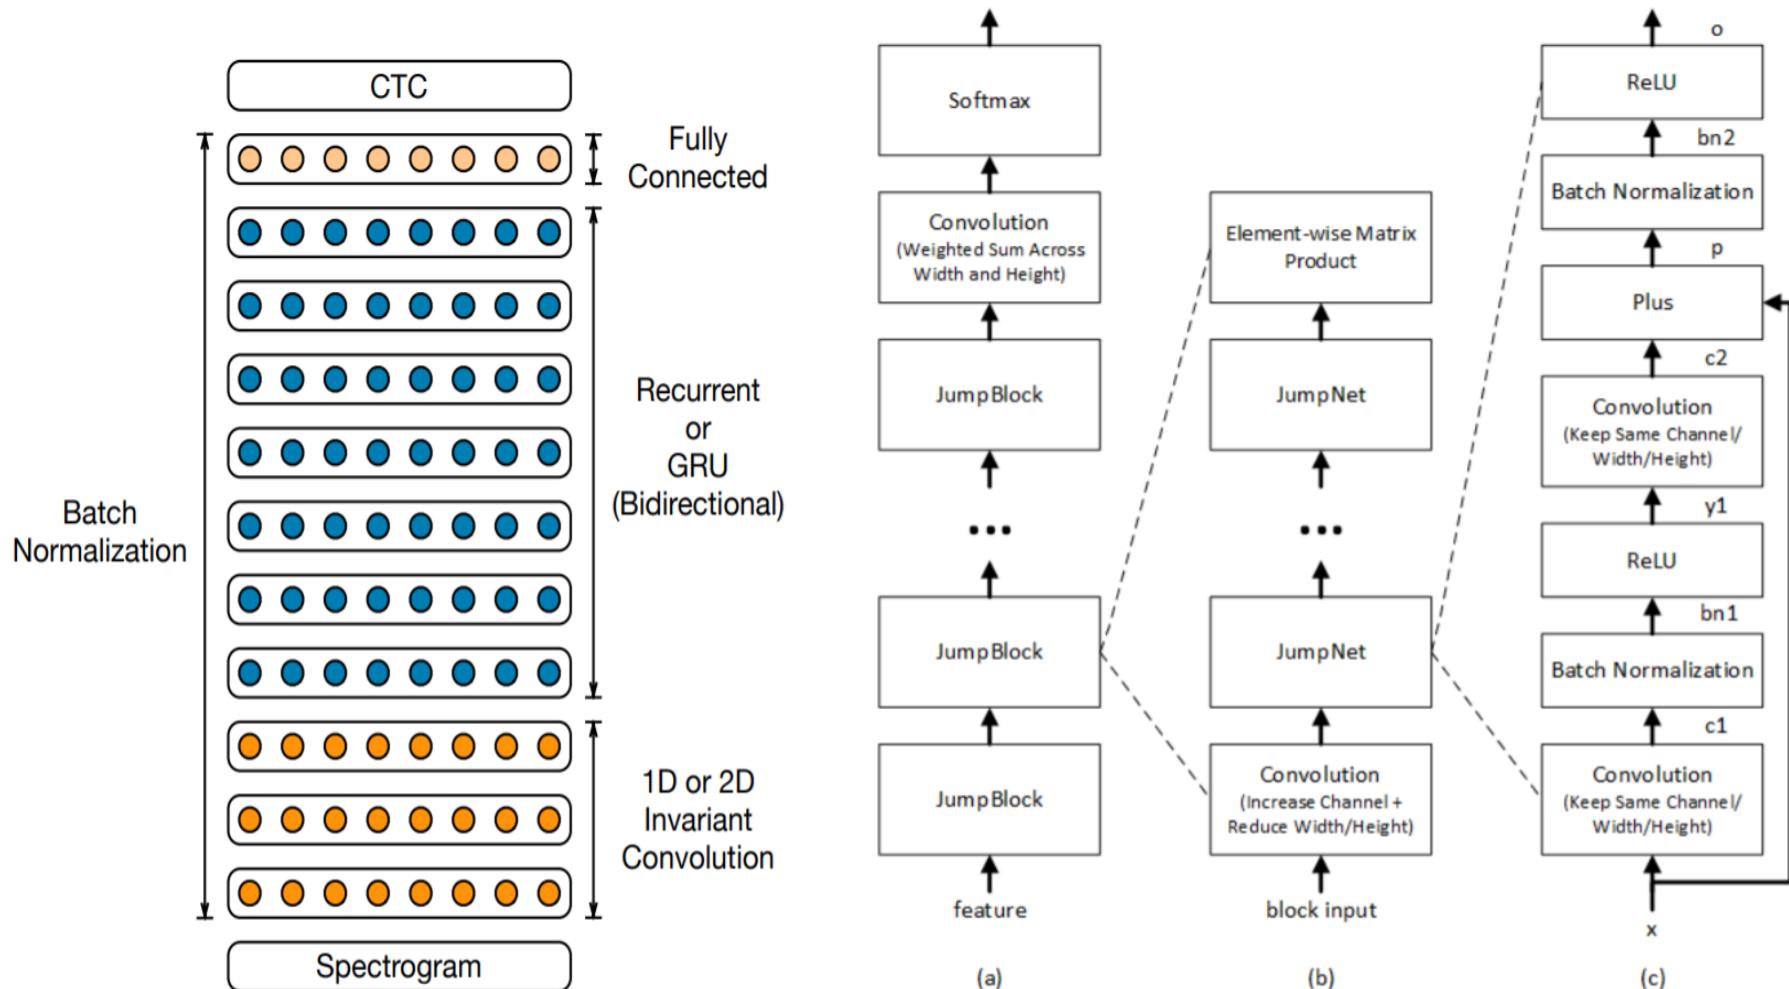
Video Super-Resolution with Motion Compensation, GCPR2017

CNN

Application for Machine Translation



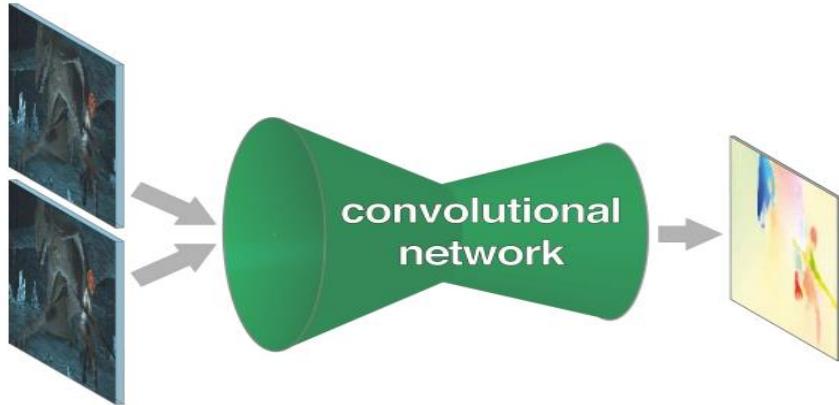
Application for Speech Recognition



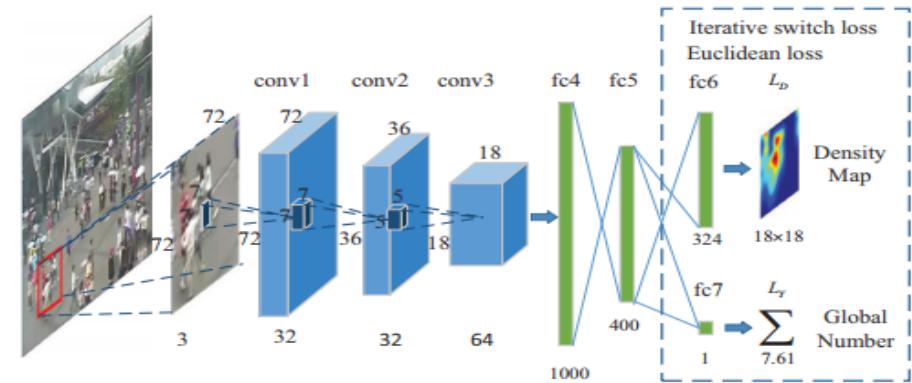
Speech Recognition in English
and Mandarin, TASLP2016

Conversational Speech Recognition, TASLP2017

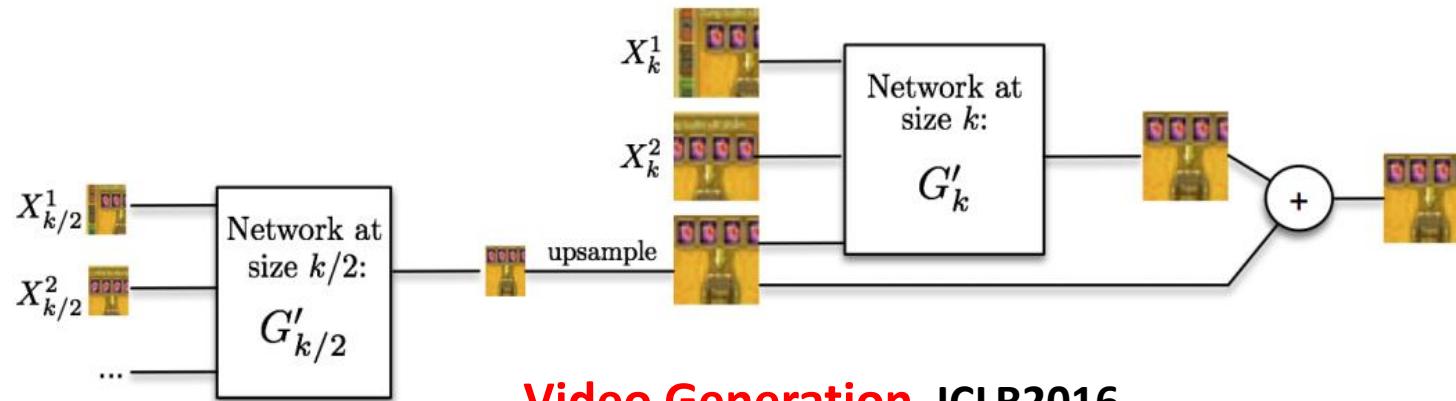
Applications for Others



Optical Flow Prediction, ICCV2015



Cross-scene Crowd Counting, CVPR2015



Video Generation, ICLR2016

Achieve state-of-the-art results in various applications

Outline

1 Course Information

2 What is Deep Learning

3 Applications of Deep Learning

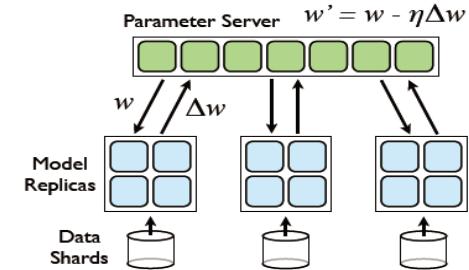
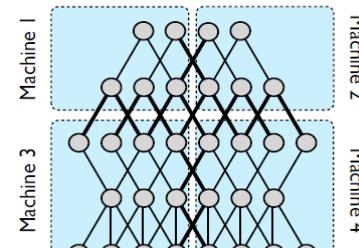
4 Future Directions

5 Project and homework guidance

Direction 1

■ Large scale deep learning

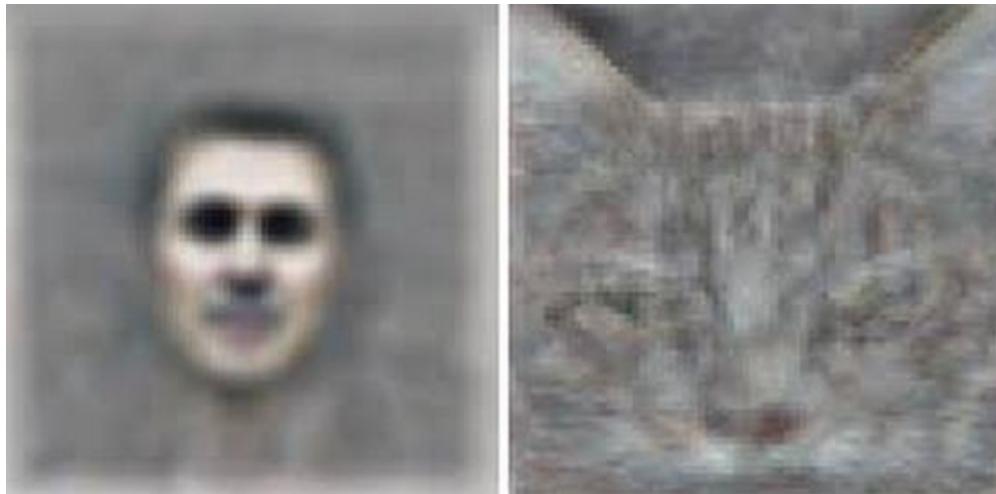
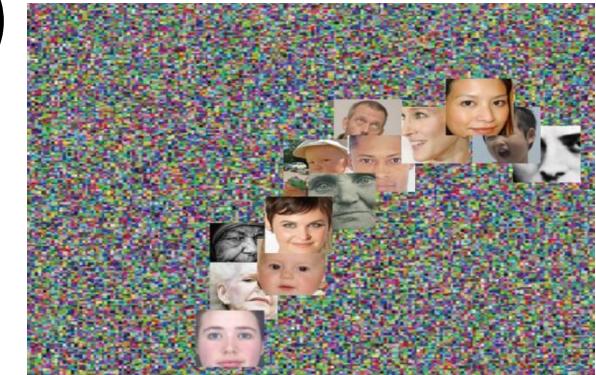
- multiple GPUs
- distributed system



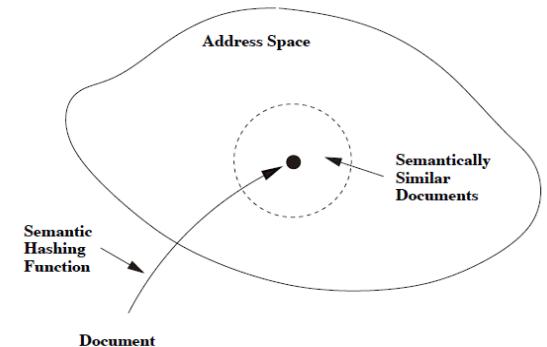
Direction 2

■ Unsupervised (semi-supervised) video feature Learning

- lack of supervised information in real applications



Unsupervised learned features



Manifold assumption

Direction 3

Multimodal learning

— videos are closely related to other data modalities, e.g., text, audio



Image classification

man woman street building
people kiss walk

Image captioning

Man kissing woman on a street

Visual QA

Q: What is the sailor doing?

A: He is kissing a nurse in a white dress.

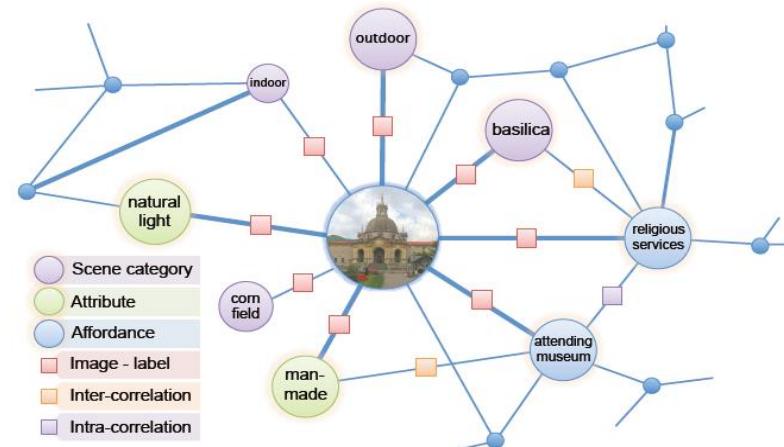
Q: Why are they kissing?

A: To celebrate V-J Day in Times Square.

Question answering



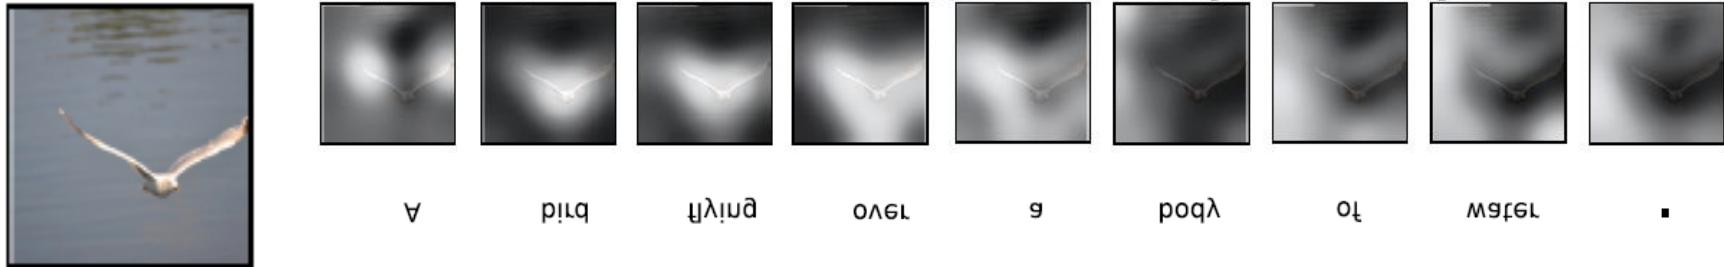
| | | |
|---|---------------------------------|--|
|  | Class auditorium | community and social work, taking class for personal interest, religious practices, waiting, attending the performing arts |
|  | Affordances landing deck | transportation and material moving work, in transit / traveling, military work |
|  | Attributes | transporting things or people, asphalt, natural light, far-away horizon, man-made |
| | | no horizon, cluttered space, dirty, eating, waiting in line |



Multimodal Knowledge Base
[Zhu et al., arXiv16]

Direction 4

- Brain-inspired deep models
 - conventional neural networks are inspired by ***human cognitive mechanism***
 - refer to advanced achievements in the fields of neuroscience and brain science



Attention-based image caption [Xu et al., ICML15]

References

1. **(Image Recognition) AlexNet** : Alex Krizhevsky, Ilya Sutskever, Geoffrey E. Hinton, ImageNet Classification with Deep Convolutional Neural Networks, NeurIPS, 2012
2. **(Action Recognition)**: Shuiwang Ji, Wei Xu, Ming Yang, Kai Yu, 3D Convolutional Neural Networks for Human Action Recognition, ICML, 2010
3. **(Detection) R-CNN**: Ross Girshick, Jeff Donahue, Trevor Darrell, Jitendra Malik, Rich feature hierarchies for accurate object detection and semantic segmentation, CVPR, 2014
4. **(Detection) Faster R-CNN**: Shaoqing Ren, Kaiming He, Ross Girshick, Jian Sun, Faster R-CNN: Towards Real-Time Object Detection with Region Proposal Networks, arXiv:1506.01497
5. **(Segmentation) F-CNN**: Jonathan Long, Evan Shelhamer, Trevor Darrell, Fully Convolutional Networks for Semantic Segmentation, CVPR, 2015
6. **(Tracking)** Chao Ma, Jia-Bin Huang, Xiaokang Yang and Ming-Hsuan Yang, Hierarchical Convolutional Features for Visual Tracking, ICCV, 2015
7. **(Super-resolution)** Chao Dong, Chen Change Loy, Kaiming He, Xiaoou Tang, Learning a Deep Convolutional Network for Image Super-Resolution, ECCV, 2014

References

8. **(Edge Detection)**: Gedas Bertasius, Jianbo Shi, Lorenzo Torresani, DeepEdge: A Multi-Scale Bifurcated Deep Network for Top-Down Contour Detection, CVPR, 2015
9. **(Face Recognition)**: Yaniv Taigman, Ming Yang, Marc'Aurelio Ranzato, Lior Wolf, DeepFace: Closing the Gap to Human-Level Performance in Face Verification, CVPR, 2014
10. **(Question Answering)**: Hyeonwoo Noh, Paul Hongsuck Seo, and Bohyung Han, Image Question Answering using Convolutional Neural Network with Dynamic Parameter Prediction, arXiv:1511.05765
11. **(Video Caption)**: Jeff Donahue, Lisa Anne Hendricks, Sergio Guadarrama, Marcus Rohrbach, Subhashini Venugopalan, Kate Saenko, Trevor Darrell, Long-term Recurrent Convolutional Networks for Visual Recognition and Description, CVPR, 2015
12. **(Text Classification)**: Xiang Zhang, Junbo Zhao, Yann LeCun, Character-level Convolutional Networks for Text Classification, NeurIPS, 2015
13. **(Retrieval)**: Fang Zhao, Yongzhen Huang, Liang Wang, Tieniu Tan, Deep Semantic Ranking Based Hashing for Multi-Label Image Retrieval, CVPR, 2015

Outline

1 Course Information

2 What is Deep Learning

3 Applications of Deep Learning

4 Future Directions

5 Project and homework guidance

1.Human Action Recognition

Project idea: Suppose that we are given a few video clips. Can we classify whether these videos contain humans ? If yes, can we track the humans in the videos, detect the pose of human body parts, and eventually infer what activity the human make? In this project, you are encouraged **to implement your own human action recognition system.** Since this area of research is so vast, we do not recommend you to try a general problem. More specified is better. Here are some references as a starting point.

Following is a list of data sets you could use.

- [1] KTH:<http://www.nada.kth.se/cvap/actions/>
- [2] Weizmann:<http://www.wisdom.weizmann.ac.il/~vision/SpaceTimeActions.html>
- [3] Hollywood Human Actions
dataset:<http://www.irisa.fr/vista/Equipe/People/Laptev/download.html>
- [4] VIRAT Video Dataset:<http://www.viratdata.org/>
(http://groups.inf.ed.ac.uk/calvin/articulated_human_pose_estimation_code/)

1.Human Action Recognition

Here are some references as a starting point.

- [1] CVPR 2011 Tutorial on Human Activity Recognition
(<http://cvrc.ece.utexas.edu/mryoo/cvpr2011tutorial/>)
- [2] Human Activity Recognition Summer course
(<http://www.cs.sfu.ca/~mori/courses/cmpt888/summer10/>)
- [2] Stanford Vision lab (http://vision.stanford.edu/discrim_rf/)
(<http://ai.stanford.edu/~bangpeng/ppmi.html>)
- [3] Poselet (<http://www.cs.berkeley.edu/~lbourdev/poselets/>)
- [4] 2D articulated human pose estimation

You could use any code from the web for computing spatial-temporal features. One good example is the spatial-temporal interest point proposed by Piotr Dollar. Source code available at <http://vision.ucsd.edu/~pdollar/research/research.html>.

2. Image Categorization

Project idea: Image categorization/object recognition has been one of the most important research problems in the computer vision community. Researchers have developed a wide spectrum of different local descriptors, feature coding schemes, and classification methods.

In this project, **you will implement your own object recognition system**. You could use any code from the web for computing image features, such as SIFT, HoG, etc.

For computing SIFT features, you could use <http://www.vlfeat.org/~vedaldi/code/sift.html>.

You're also encouraged to try the state-of-the-art deep learning methods:

<http://deeplearning.net/>

Following is a list of data sets you could use.

[1] Caltech101/256:

http://www.vision.caltech.edu/Image_Datasets/Caltech101/Caltech101.html

[2] The PASCAL Object Recognition Database Collection:

<http://pascallin.ecs.soton.ac.uk/challenges/VOC/databases.html>

[3] LabelMe:<http://labelme.csail.mit.edu/>

[4] Face in the wild:<http://vis-www.cs.umass.edu/lfw/>

[5] ImageNet:<http://www.image-net.org/index>

[6] TinyImage:<http://groups.csail.mit.edu/vision/TinyImages/>

3. Multi-modal Social Media

Project idea: In this project, we encourage students to infer the underlying relations between different modalities of information on the Web. Here are some examples.

(1) Given a photo of movie poster (image), can we retrieve related trailers (video) or latest news articles (documents) of the movie?

To make project simpler, we recommend focusing on less than five movies (eg. 'Rise of the planet of the apes' and 'The smurfs'). You first download posters and trailers from some well-organized sites such as imdb.com or itunes.com. They will be used as training data to learn your classifiers. Now your job is to gather raw data from youtube.com or Flickr, and classify them. In this project, we encourage you to explore the possibility to build classifiers to be learned from one information modality (eg. images), and to be applicable to other modalities (eg. trailer videos).

3. Multi-modal Social Media

(2) Given a beer label (image), can we search for which frames of a given video clip the logo or bottle appears?

Suppose that you are a big fan of Guinness beer. You can easily download the clean Guinness logo or cup images by Google image search. These images can be used to learn your detector, which can discover the frames that the logo appears in the video clips. For testing, you can download some video clips from youtube.com.

The above examples are just two possible candidates, and any new ideas or problem definitions are welcome.

For this purpose, one may take advantage of some source codes available on the Web as unit modules (e.g., near-duplicated image detection, object recognition, action recognition in video).

Another interesting direction is to improve the current state-of-the-arts methods by considering more practical scenarios.

3. Multi-modal Social Media

Related Papers and Software::

- A good example of how a machine learning technique is successfully applied to real systems (ex. Google news recommendation).
- [1] Das, Datar, Garg, Rajaram. Google news personalization: scalable online collaborative filtering. WWW 2007.
- One of most popular approaches to near duplicated image detection is LSH families.
- [2] <http://www.mit.edu/~andoni/LSH/> (This webpage links several introductory articles and source codes).
- Various hashing techniques in computer vision (papers and source codes).
- [3] Spectral Hashing (<http://www.cs.huji.ac.il/~yweiss/SpectralHashing/>)
- [4] Kernelized LSH (<http://www.eecs.berkeley.edu/~kulis/klsh/klsh.htm>)
- Recognition in video
- [5] Naming of Characters in Video (<http://www.robots.ox.ac.uk/~vgg/data/nface/index.html>)
- [6] Action recognition in Video
(<http://www.robots.ox.ac.uk/~vgg/data/stickmen/index.html>)
- Recognition in images
- [7] Human pose detection (Poselet) (<http://www.eecs.berkeley.edu/~lbourdev/poselets/>)
- [8] General object detection (<http://people.cs.uchicago.edu/~pff/latent/>)

Other Projects

4. Image Segmentation
5. Face Recognition
6. Text Classification
7. Question Answering
8. Image Denosing/Super-resolution
9. Image Retrieval
10. Tracking
11.

Note that not all topics are equally difficult. But we'll take this into consideration when evaluating each team's performance and make the assessment as fair as possible.

Homework

- Homework will be assigned **after every 2-3 classes**, you should submit the reading report **in 2 weeks**;
- We will provide some research directions and corresponding paper list, each student need to choose one paper and submit a reading report **no more than 2 pages**;
- You should briefly explain the **problem** the paper is working on, and the proposed **solutions**;
- Your **own understanding of the contributions** and **unsolved problems** of the paper are welcomed in the reading report;

Acknowledgement

Some of the materials in these slides are drawn inspiration from:

- Alexander Amini, MIT University, Introduction to Deep Learning course
- Hung-yi Lee, National Taiwan University, Machine Learning and having it Deep and Structured course
- Chenglin Liu, CASIA, Pattern Recognition course
- Fei-Fei Li, Standord University, CS231n Convolutional Neural Networks for Visual Recognition course

Questions?

Thank You !



WeChat Group for Deep Learning